

ECE7121 Learning-based control – 2025 Fall

Environment and Simulation



INHA UNIVERSITY

Overview

- > IL -> RL / control problem
- > Multi-arm bandit
- > MDP / POMDP
- > Continuity
- > Dynamic system model
- > Simulation
- > Simulators

Limitations of imitation learning

- > We need experts (e.g., human)
 - Could be expensive
- > Can not go beyond the expert level
 - We want super intelligence
- > Expert demo could be multi-modal and not optimal
 - Learning might go wrong

RL Success

> Game




AlphaGo: Go World champion



AlphaStar: Grandmaster (99.8%)


RL = trial-and-error learning

- > Reinforcement in educational psychology
 - Big advantage: AI can learn autonomously



RL = trial-and-error learning

- > Learning a policy that maximizes rewards by interacting with the environment



- Each action results in an immediate reward.
- We want to choose actions that maximize our immediate reward in expectation.

Multi-armed bandits

> bandit =



One-armed bandit: slot machine



Multi-armed bandit: multiple slot machines

Multi-armed bandits

- > The state does not change!
 - We don't move. We use the same slot machines.
- > We have N slot machines and select one to pull.
 - We have N possible actions.
- > We will get an immediate reward from the pulled slot machine.
 - Each slot machine gives a random reward.
 - The reward probability of each machine is fixed but unknown.
- > Objective: maximize cumulative rewards

How to earn money?

- > Strategy 1: pull each once, exploit the best



Reward: 100₩

0₩

200₩

30₩

- > Next, we only pull the third machine

How to earn money?

- > Strategy 2: pull each 4 times, exploit the best



Reward:

100₩	0₩	200₩	30₩
100₩	300₩	0₩	500₩
100₩	0₩	0₩	20₩
100₩	400₩	0₩	40₩
400₩	700₩	200₩	590₩

- > Next, we only pull the second machine

Achieving a balance


- > Pulling the same machine several times
 - = learning reward probability distribution and its mean
 - = **exploration (collecting information)**
- > Pulling the best machine
 - = exploiting the best to earn money (given current information)
 - = **exploitation (collecting reward)**
- > Action-value for action a is its mean reward
 - $Q_t^*(a) = \mathbb{E}[R_{t+1} | A_t = a]$, action-value estimate: $Q_t(a) \approx Q_t^*(a)$
 - $A_t^* = \arg \max Q_t(a)$
 - If $A_t = A_t^*$: exploiting
 - If $A_t \neq A_t^*$: exploring
- > We need to do both

Exploration dilemma

- > Exploration vs Exploitation dilemma
 - The best long-term strategy may involve short-term sacrifices
 - This is not a problem unique to RL; it is a fundamental issue in the decision making of any intelligent agent.
- > Restaurant selection
 - exploitation: go to your favorite restaurant
 - exploration: try a new restaurant
- > Studying
 - exploitation: solve example problems
 - exploration: read additional materials

Exploration dilemma

- > ϵ – greedy algorithm





Markov decision process

- > Multi-armed bandit
 - $\tau: (A_t, R_{t+1}, A_{t+1}, R_{t+2}, A_{t+2}, R_{t+3}, \dots)$
- > Markov decision process
 - $\tau: (S_t, A_t, R_{t+1}, S_{t+1}, A_{t+1}, R_{t+2}, S_{t+2}, A_{t+2}, R_{t+3}, \dots)$
- > Markov property
 - $p[R_{t+1} = r, S_{t+1} = s' | S_0, A_0, R_0, S_1, A_1, R_1, \dots, S_t, A_t]$
 $= p[R_{t+1} = r, S_{t+1} = s' | S_t, A_t]$
 - Only the present determines the future; we can ignore the history.

Markov decision process

- > Finite Markov decision process is a tuple (S, A, T, r, γ)
 - S is a finite set of states $s \in S$
 - A is a finite set of actions $a \in A$
 - T is one step transition/dynamics function $p(s'|s, a)$
 - $r(s, a, s')$ is a reward function
 - γ is a discount factor ($0 \leq \gamma \leq 1$)
- > $\pi(a|s)$: a policy is a distribution over actions given states




Markov decision process

> Maximize your sum of future rewards (cumulative reward)


- $G = R(\tau) = R_1 + R_2 + R_3 + R_4 \dots$
- Future rewards may less important
- $G = R(\tau) = R_1 + \gamma R_2 + \gamma^2 R_3 + \gamma^3 R_4 + \dots$

> The robot collects toll on every bridge




POMDP

- > Partially Observable MDP
 - Finite POMDP is a tuple $(S, A, O, T, h, r, \gamma)$
 - h is the observation model $p(o|s)$
- > Learn a policy $\pi(a|o)$



POMDP

> Playing Atari with DQN



RL in robot control

- > Now, we want to control diverse robotic platforms
 - Can we formulate a control problem mathematically?






- > Optimal control = optimization of dynamic systems

Optimal control

- > Optimal control problem (OCP) have three ingredients:
 - dynamic system model
 - constraints
 - objective function
- > Typical control methods for complex robotic systems:
 - 1. Model the system
 - 2. Define a desired behavior (cumulative reward)
 - 3. Find the trajectory that maximizes the goal
 - 4. Find the inverse model
 - 5. Find the actuation to follow

Continuity

		State space	
		Discrete	Continuous
Time space	Continuous		
	Discrete		

Dynamic system models

- > Discrete time

- $s_{k+1} = f(s_k, a_k)$

- > Continuous time


- Ordinary differential equation (ODE)
 - $\dot{s} = \frac{ds}{dt}(t) = f_c(s(t), a(t))$

- > Stochastic

- Adding random disturbance ϵ

Control problem to MDP

- > Transform a continuous system into a discrete system
 - define $s_k = s(t_k)$ on equidistant time grid $t_k = k\Delta t$ with sampling time
 - use zero order hold control $a(t) = a_k$ on $t \in [t_k, t_{k+1}]$
 - use numerical simulation to compute ODE solution
(Euler integration, Runge-Kutta RK4)





Numerical simulation and integration

- > Simplest implementation is a single step of an Euler integrator
 - $s' = f(s, a) = s + \Delta t f_c(s, a)$
 - $\Delta t \rightarrow 0$, simulation becomes more accurate
- > Even if we have a fixed time step Δt ,
 - more accurate simulation can be made by iterating an Euler integrator
 - for $i=1:N$, $x_{i+1} = \frac{\Delta t}{N} f_c(s, a)$
- > Runge Kutta method is more accurate and efficient
 - Runge Kutta 4th order
$$f(s, a) = s + \Delta t / 6 (v_1 + 2v_2 + 2v_3 + v_4)$$
$$v_1 := f_c(s, a)$$
$$v_2 := f_c(s + (\Delta t / 2) v_1, a)$$
$$v_3 := f_c(s + (\Delta t / 2) v_2, a)$$
$$v_4 := f_c(s + \Delta t v_3, a)$$

Numerical simulation and integration


- > $\dot{s} = s + a$, for $\Delta t = 1$. Exact solution is $f(s, a) = e = 2.718$
- > Four Euler steps give 2.441 (10.2% error), RK4 gives 2.708 (0.36% error)



Dynamics model example


> Bicycle model

$$\frac{d}{dt} \begin{bmatrix} x \\ y \\ \psi \\ v \end{bmatrix} = \begin{bmatrix} v \cos(\psi + \beta) \\ v \sin(\psi + \beta) \\ \frac{v}{L_r} \sin \beta \\ a \end{bmatrix}, \text{ with } \beta := \arctan\left(\frac{L_r}{L_r + L_f} \arctan \delta\right)$$



Dynamics model example

> Drone dynamics




$$\begin{cases} \dot{\mathbf{p}} = \mathbf{v} \\ \dot{\mathbf{v}} = \mathbf{g} + \frac{1}{m} \mathbf{R} \mathbf{f}_t \\ \mathbf{J} \dot{\boldsymbol{\omega}} = \boldsymbol{\tau} - \boldsymbol{\omega} \times \mathbf{J} \boldsymbol{\omega} \end{cases}$$

$$\begin{bmatrix} \dot{p}_x \\ \dot{p}_y \\ \dot{p}_z \\ \dot{v}_x \\ \dot{v}_y \\ \dot{v}_z \\ \dot{\phi} \\ \dot{\theta} \\ \ddot{\psi} \\ \dot{\phi} \\ \ddot{\theta} \\ \dot{r} \\ \dot{\mathbf{a}}_t \end{bmatrix} = \begin{bmatrix} v_x \\ v_y \\ v_z \\ \dot{v}_x^{(*)} \\ \dot{v}_y^{(*)} \\ \dot{v}_z^{(*)} \\ \dot{\phi} \\ \dot{\theta} \\ \frac{t_\phi}{c_\theta} \dot{\theta} + \frac{1}{c_\phi c_\theta} r \\ -\omega_{n,\phi}^2 \phi - 2\zeta_\phi \omega_{n,\phi} \dot{\phi} + \omega_{n,\phi}^2 \bar{\phi}_r \\ -\omega_{n,\theta}^2 \theta - 2\zeta_\theta \omega_{n,\theta} \dot{\theta} + \omega_{n,\theta}^2 \bar{\theta}_r \\ -\sigma_r r + \sigma_r \bar{r}_r \\ -\sigma_t \mathbf{a}_t + K \left(\frac{U}{U_n} \right)^\alpha \sigma_t \bar{T} \end{bmatrix}$$

$$(*) \underbrace{\begin{bmatrix} \dot{v}_x \\ \dot{v}_y \\ \dot{v}_z \end{bmatrix}}_{\dot{\mathbf{v}}} = \underbrace{\mathbf{R} \begin{bmatrix} 0 \\ 0 \\ \mathbf{a}_t \end{bmatrix}}_{\mathbf{a}_t} - \underbrace{\mathbf{R} \mathbf{D} \mathbf{R}^\top}_{\mathbf{a}_v} \begin{bmatrix} v_x \\ v_y \\ v_z \end{bmatrix} - \underbrace{\begin{bmatrix} 0 \\ 0 \\ g \end{bmatrix}}_g$$

General stochastic models

- > We always have some random noise ϵ_k
 - disturbances or measurement errors



State space and MDP

> MDP state space view

- $s_{k+1} = f(s_k, a_k, \epsilon_k)$
- $y_k = s_k$

> MDP probabilistic view

- $p(s_{k+1}|s_k, a_k)$
- $p(y_k|s_k, a_k) = \delta_d(y_k - s_k)$

> POMDP state space view

- $s_{k+1} = f(s_k, a_k, \epsilon_k)$
- $y_k = g(s_k, a_k, \epsilon_k)$

> POMDP probabilistic view

- $p(s_{k+1}|s_k, a_k)$
- $p(y_k|s_k, a_k)$


Robot simulators

- > Computational physics and computer graphics
- > Simulator taxonomy by simulately.wiki

Simulator	Physics Engine	Rendering	Sensor 	Dynamics	GPU-accelerated Simulation	Open-Source
IsaacSim	PhysX 5	Rasterization; RayTracing	RGBD; Lidar; Force; Effort; IMU; Contact; Proximity	Rigid;Soft;Cloth;Fluid	✓	✗
IsaacGym	PhysX 5, Flex	Rasterization;	RGBD; Force; Contact;	Rigid;Soft;Cloth	✓	✗
SAPIEN	PhysX 5, Warp	Rasterization; RayTracing 	RGBD; Force; Contact;	Rigid;Soft;Fluid	✓	✓
Pybullet	Bullet	Rasterization;	RGBD; Force; IMU; Tactile;	Rigid;Soft;Cloth	✗	✓
MuJoCo	MuJoCo	Rasterization;	RGBD; Force; IMU; Tactile;	Rigid;Soft;Cloth	✓ 	✓
CoppeliaSim	MuJoCo; Bullet; ODE; Newton; Vortex	Rasterization; RayTracing 	RGBD; Force; Contact;	Rigid;Soft;Cloth	✗	✓
Gazebo	Bullet; ODE; DART; Simbody	Rasterization;	RGBD; Lidar; Force; IMU;	Rigid;Soft;Cloth	✗	✓
Genesis	Genesis	Rasterization; Raytracing	RGBD; Tactile; (update soon)	Rigid;Soft;Cloth;Fluid	✓	✓


Gazebo

- > Developed from 2002
- > Open-source robotic simulator
 - Working well with ROS
 - GPU acceleration is not supported
 - Physics engine: Bullet, ODE, DART, Simbody




PyBullet

- > Python + Bullet physics engine
- > Free and open source, easy to use
 - GPU acceleration is not supported




MuJoCo

- > Developed by Roboti LLC, acquired and made freely available by DeepMind
 - Multi-Joint Dynamics-with-Contact
 - Accurate and efficient contact dynamics (compliant contact model)
 - Differentiable



RaiSim

- > Developed by Jemin Hwangbo (KAIST)
 - Four major computational components in rigid-body simulators
 - Collision detection, forward dynamics computation, contact properties computation, contact solver
 - Visualization options – Unity, Unreal, Ogre




IssacSim

- > Developed by NVIDIA
 - End-to-end GPU accelerated
 - Massive parallelization of thousands of environments
 - PhysX 5 physics engine, including fluid dynamics




Why simulators for robot learning

- > RL is very sample inefficient



Why simulators for robot learning

- > Google QT-Opt
 - 4 months, 800 robot hours, 7 robots, 580,000 attempts




Why simulators for robot learning

- > Advantages of using simulated data
 - Cheap, fast, and scalable
 - Safe
 - Labeled (we have access to ground truth)
 - No physical harm or deformation (wear and tear)

- > Disadvantages
 - Not exactly same with real env. (sim2real)



Aerodynamics in wind, *Neural-Fly*



Fluid dynamics, MIT van Rees Lab




Offroad vehicle dynamics, UW Racer team

Environment for RL

> Gymnasium

- Provide an API for all single agent RL envs.
- Maintained by OpenAI




```
env = gym.make  
observation, info = env.reset  
action = env.action_space.sample()  
observation, reward, terminated,  
truncated, info = env.step(action)
```

Exploiting abundant data


- > RL from wild animal videos



Step 1: Reward Learning From Wild Animal Videos



Step 2: Physical Grounding with RL in a Simulator



More on simulators

- > <https://github.com/knm McGuire/best-of-robot-simulators>
- > Gymnasium: <https://gymnasium.farama.org/>
- > Gymnasium-robotics: <https://robotics.farama.org/>
- > 2D sim: <https://ir-sim.readthedocs.io/en/stable/>
- > Drone: <https://github.com/utiasDSL/gym-pybullet-drones>
- > Autonomous driving: <https://carla.org/>
- > Manipulator / Legged robot / Humanoid (contact-dynamics):
MuJoCo, Issac Sim, RaiSim, Genesis ...