

ECE7121 Learning-based control – 2025 Fall

# RL basics (Classical RL)



INHA UNIVERSITY


# Overview

- > Tabular-value based RL (Classic RL)
- > Value function
- > Policy evaluation
- > Policy improvement
  - policy iteration
  - value iteration
- > Temporal difference learning
- > Q-learning

# Major components of an RL agent


- > An RL agent may include those components:
  - policy: agent's behavior function
    - deterministic policy  $A_t = \pi(s)$
    - stochastic policy  $A_t = \pi(a|s) = p(A_t = a|S_t = s)$
  - value function: how good is each state and/or action
  - model: agent's representation of the environment
    - predict what the environments will do next
    - next state:  $p(S_{t+1} = s'|S_t = s, A_t = a)$
    - next reward:  $p(R_{t+1}|S_t = s, A_t = a)$

## Maze example




## Maze example

### > Policy




## Maze example

### > Value function



# Maze example



## > Model



- Dynamics: how actions change the state
- Rewards: how much reward from each state
- Grid layout represents transition model
- Numbers represent immediate reward

## The goal of RL

- > Finite horizon case: T is finite
- > Infinite horizon case: T =  $\infty$
- > The cumulative reward is often discounted:  $G_t = \sum_t \gamma^t r(s_t, a_t)$
- > Goal: find a policy to maximize the cumulative reward



## Value functions

- > State-value function ( $V$ ): the expected return starting from state  $s$ , and then following policy  $\pi$ 
  - $V_\pi = \mathbb{E}_\pi[G_t | S_t = s]$
  - optimal:  $V^*(s) = \max_\pi V_\pi(s)$   
(maximum value function over all policies)  
(the expected return when acting optimally)
- > Action-value function ( $Q$ ): the expected return starting from state  $s$ , taking action  $a$ , and then following policy  $\pi$ 
  - $Q_\pi(s, a) = \mathbb{E}_\pi[G_t | S_t = s, A_t = a]$
  - optimal:  $Q^*(s, a) = \max_\pi Q_\pi(s, a)$
- > Value functions capture the knowledge of the agent regarding how good is each state for the goal the agent is trying to achieve.

## Value functions


- > Once we obtain the optimal value function, we can find an optimal policy
- > Maximizing over  $Q^*(s, a)$ 
  - $\pi^*(a|s) = \begin{cases} 1 & \text{if } a = \operatorname{argmax}_a Q^*(s, a) \\ 0 & \text{otherwise} \end{cases}$
- > Maximizing over  $V^*(s)$  with the model dynamics
  - $\pi^*(a|s) = \begin{cases} 1 & \text{if } a = \operatorname{argmax}_a \sum p(s', r|s, a)(r + \gamma V^*(s')) \\ 0 & \text{otherwise} \end{cases}$
  - We need the dynamics to do one step lookahead to choose the optimal  $a$

## Roadmap

- > Computing state and state-action value functions by solving linear systems of equations
- > The matrix inversion is too costly
  - > iterative estimation is required (Bellman backup operation)
- > We cannot visit every state
  - > selective backups on state-actions that the agent visits
- > We may not know dynamics
  - > Monte Carlo learning or TD learning

## Recursive relationships for returns

$$\begin{aligned}
 > G_t &= R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \gamma^3 R_{t+4} + \dots \\
 &= R_{t+1} + \gamma(R_{t+2} + \gamma R_{t+3} + \gamma^2 R_{t+4} + \dots) \\
 &= R_{t+1} + \gamma G_{t+1}
 \end{aligned}$$



> By taking expectations

- $\mathbb{E}[G_t | S_t = s] = \mathbb{E}[R_{t+1} + \gamma G_{t+1} | S_t = s]$
- $V_\pi(s) = \mathbb{E}[R_{t+1} + \gamma V_\pi(s') | S_t = s]$   
 $= \sum_{s'} p(s', r | s)[r + \gamma V_\pi(s')]$  (Bellman expectation equation)

- For all states,  $V_\pi = R_\pi + \gamma P_\pi V_\pi$   
where  $v$  is a column vector with one entry per state

$$\begin{bmatrix} V(1) \\ \vdots \\ V(n) \end{bmatrix} = \begin{bmatrix} R_1 \\ \vdots \\ R_n \end{bmatrix} + \gamma \begin{bmatrix} P_{11} & \cdots & P_{1n} \\ \vdots & \ddots & \vdots \\ P_{n1} & \cdots & P_{nn} \end{bmatrix} \begin{bmatrix} V(1) \\ \vdots \\ V(n) \end{bmatrix}$$



## Solving the Bellman expectation equation

- > The Bellman expectation equation is a linear equation
- > It can be solved directly:
  - $$V_\pi = R_\pi + \gamma P_\pi V_\pi$$
$$(I - \gamma P_\pi)V_\pi = R_\pi$$
$$V_\pi = (I - \gamma P_\pi)^{-1}R_\pi$$
- > Computational complexity is  $O(n^3)$  for  $n$  states
- > There are many iterative methods for large state system
  - Dynamic programming
  - Monte-Carlo evaluation
  - Temporal-Difference learning

## Relating state and state-action value functions


- > The action-value function can similarly be decomposed

$$Q_\pi(s, a) = \mathbb{E}[R_{t+1} + \gamma Q_\pi(s', a') | S_t = s, A_t = a]$$



$$V_\pi(s) = \sum_a \pi(a|s) Q_\pi(s, a)$$

$$Q_\pi(s, a) = \sum_{s'} p(s', r|s, a) [r + \gamma V_\pi(s')]$$



$$V_\pi(s) = \sum_a \pi(a|s) \sum_{s'} p(s', r|s, a) [r + \gamma V_\pi(s')]$$

$$Q_\pi(s, a) = \sum_{s'} p(s', r|s, a) [r + \gamma \sum_a \pi(a'|s') Q_\pi(s', a')]$$

## Bellman optimality equations

> For the Bellman expectation equations, we sum over all the possibilities.

> Now, we choose only the best action.

- $V_\pi(s) = \sum_a \pi(a|s)Q_\pi(s, a) \rightarrow V^* = \max_a Q^*(s, a)$

- $Q_\pi(s, a) = \sum_{s'} p(s', r|s, a)[r + \gamma V_\pi(s')]$

$$\rightarrow Q^*(s, a) = \sum_{s'} p(s', r|s, a)[r + \gamma V^*(s')]$$

- $V_\pi(s) = \sum_a \pi(a|s) \sum_{s'} p(s', r|s, a)[r + \gamma V_\pi(s')]$

$$\rightarrow V_\pi^*(s) = \max_a \sum_{s'} p(s', r|s, a)[r + \gamma V_\pi^*(s')]$$

- $Q_\pi(s, a) = \sum_{s'} p(s', r|s, a)[r + \gamma \sum_a \pi(a'|s')Q_\pi(s', a')]$

$$\rightarrow Q^*(s, a) = \sum_{s'} p(s', r|s, a)[r + \gamma \max_a Q^*(s', a')]$$

## Solving the Bellman optimality equation

- > Bellman optimality equation is nonlinear
- > No closed form solution (in general)
- > Many iterative solution methods
  - Using models / dynamic programming
    - Value iteration
    - Policy iteration
  - Using samples
    - Monte Carlo
    - Q-learning
    - SARSA

## Extensions to MDPs

- > Infinite and continuous MDP
- > Continuous state and/or action spaces
  - closed form for linear quadratic model (LQR)
- > Continuous time
  - requires partial differential equations
  - Hamilton-Jacobi-Bellman (HJB) equation

## Planning by dynamic programming


- > Dynamic programming assumes full knowledge of the MDP
- > It is used for planning in an MDP
- > For prediction:
  - input: MDP and policy
  - output: value function  $V_\pi$
- > For control:
  - input: MDP
  - output: optimal value function  $V^*$  and optimal policy  $\pi^*$

## Policy Evaluation

- > Problem: evaluate a given policy  $\pi$  (prediction)
- > Solution: iterative application of Bellman expectation backup
  - $V_1 \rightarrow V_2 \rightarrow \dots \rightarrow V_\pi$
  - Synchronous backups – for all states
- > Iterative policy evaluation
- > 
$$V_{k+1}(s) = \sum_a \pi(a|s) \sum_{s'} p(s', r|s, a) [r + \gamma V_k(s')]$$

## Policy Evaluation

- > Evaluating a random policy in the small GridWorld



- undiscounted episodic MDP ( $\gamma = 1$ )
- one terminal state (shown twice as shaded squares)
- actions leading out of the grid leave state unchanged
- reward is -1 until the terminal state is reached
- agent follows uniform random policy  
(each action has a probability of 0.25)

# Policy Evaluation

$$> V_{k+1}(s) = \sum_a \pi(a|s) \sum_{s'} p(s', r|s, a) [r + \gamma V_k(s')]$$

$v_k$  for the  
Random Policy

$k = 0$

0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0

$k = 3$

0.0	-2.4	-2.9	-3.0
-2.4	-2.9	-3.0	-2.9
-2.9	-3.0	-2.9	-2.4
-3.0	-2.9	-2.4	0.0

$k = 1$

0.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	0.0

$k = 10$

0.0	-6.1	-8.4	-9.0
-6.1	-7.7	-8.4	-8.4
-8.4	-8.4	-7.7	-6.1
-9.0	-8.4	-6.1	0.0

$k = 2$

0.0	-1.7	-2.0	-2.0
-1.7	-2.0	-2.0	-2.0
-2.0	-2.0	-2.0	-1.7
-2.0	-2.0	-1.7	0.0

$k = \infty$

0.0	-14.	-20.	-22.
-14.	-18.	-20.	-20.
-20.	-20.	-18.	-14.
-22.	-20.	-14.	0.0

# Policy Evaluation

## > Overall algorithm

Input  $\pi$ , the policy to be evaluated

Initialize an array  $V(s) = 0$ , for all  $s \in \mathcal{S}^+$

Repeat

$\Delta \leftarrow 0$

    For each  $s \in \mathcal{S}$ :

$v \leftarrow V(s)$

$V(s) \leftarrow \sum_a \pi(a|s) \sum_{s',r} p(s',r|s,a) [r + \gamma V(s')]$

$\Delta \leftarrow \max(\Delta, |v - V(s)|)$

    until  $\Delta < \theta$  (a small positive number)

Output  $V \approx v_\pi$


## > It will converge to the fixed value function

## Policy Evaluation

- > Once we found the converged (optimal) value function, we can get a better (optimal) policy

$k = \infty$

0.0	-14.	-20.	-22.
-14.	-18.	-20.	-20.
-20.	-20.	-18.	-14.
-22.	-20.	-14.	0.0



## Policy Evaluation

- > Convergence proof
  - Define the Bellman expectation backup operator
  - $T_\pi(V) = R_\pi + \gamma P_\pi V$
  - This operator is a  $\gamma$ -contraction; it makes value functions closer by at least  $\gamma$
  - $$\begin{aligned} \|T_\pi(U) - T_\pi(V)\|_\infty &= \|R_\pi + \gamma P_\pi U - (R_\pi + \gamma P_\pi V)\|_\infty \\ &= \|\gamma P_\pi(U - V)\|_\infty \\ &\leq \gamma \|P_\pi\|_\infty \|U - V\|_\infty \\ &= \gamma \|U - V\|_\infty \end{aligned}$$
  - $T_\pi$  converges to a unique fixed point at a linear convergence rate  $\gamma$

## Policy iteration

- > Now, we want to move to the control problem
- > Policy iteration
  - Evaluate the policy
  - Improve the policy by acting greedily with respect to  $V_\pi$

$$\pi' = \text{greedy}(V_\pi)$$

- It always converges to optimal policy  $\pi^*$


# Policy iteration

$v_k$  for the Random Policy

0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0


$k = 0$

Greedy Policy w.r.t.  $v_k$




$k = 1$

0.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	0.0



$k = 2$



0.0	-1.0	-2.0	-2.0
-1.0	-2.0	-2.0	-2.0
-2.0	-2.0	-2.0	-1.0
-2.0	-2.0	-1.0	0.0



$k = 3$

0.0	-1.0	-2.0	-3.0
-1.0	-2.0	-3.0	-2.0
-2.0	-3.0	-2.0	-1.0
-3.0	-2.0	-1.0	0.0

## Policy iteration



## Policy iteration

- > Consider a deterministic policy,  $a = \pi(s)$
- > Improve the policy by acting greedily,  $\pi'(s) = \operatorname{argmax}_a Q_\pi(s, a)$
- > It improves the value function

- $$Q_\pi(s, \pi'(s)) = \max_a Q_\pi(s, a) \geq Q_\pi(s, \pi(s)) = V_\pi(s)$$
- $$\begin{aligned} V_\pi(s) &\leq Q_\pi(s, \pi'(s)) = \mathbb{E}_{\pi'}[R_{t+1} + \gamma V_\pi(s') | S_t = s] \\ &\leq \mathbb{E}_{\pi'}[R_{t+1} + \gamma Q_\pi(s', \pi'(s')) | S_t = s] \\ &= \mathbb{E}_{\pi'}[R_{t+1} + \gamma \mathbb{E}_{\pi'}[R_{t+2} + \gamma V_\pi(s'')] | S_t = s] \\ &\leq \mathbb{E}_{\pi'}[R_{t+1} + \gamma R_{t+2} + \gamma^2 Q_\pi(s', \pi'(s'')) | S_t = s] \\ &\quad \dots \\ &\leq \mathbb{E}_{\pi'}[R_{t+1} + \gamma R_{t+2} + \gamma R_{t+3} + \dots | S_t = s] = V_{\pi'}(s) \end{aligned}$$

## Policy iteration

- > If improvements stop,  $Q_\pi(s, \pi'(s)) = V_\pi(s)$
- > Then, the Bellman optimality equation has been satisfied

$$V^* = \max_a Q^*(s, a)$$


- > Therefore,  $V$  and  $\pi$  are optimal

## Generalized policy iteration

- > Any interleaving of policy evaluation and policy improvement
  - > Evaluation and improvement do not need to be exact or complete at each step
    - partial updates, approximations, or asynchronous updates are allowed
  - > All RL methods are a form of GPI
- 
- > Policy iteration: full evaluation + full improvement
    - using Bellman expectation eqn.
  - > Value iteration: one-step evaluation
    - using Bellman optimality eqn.
- 
- > How GPI leads optimality?
    - Over time, even with approximate steps the policy becomes progressively better

## Value iteration

- > Problem: find optimal policy  $\pi$
- > Solution: iterative application of Bellman optimality backup
- > Using synchronous backups
- > Unlike policy iteration, there is no explicit policy
- > 
$$V_{k+1}(s) = \max_a [R(s, a) + \gamma \sum_{s'} p(s', r | s, a) V_k(s')]$$



## Value iteration

$$> V_{k+1}(s) = \max_a [R(s, a) + \gamma \sum_{s'} p(s', r|s, a)V_k(s')]$$

<b>g</b>				

Problem

0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

$V_1$

0	-1	-1	-1
-1	-1	-1	-1
-1	-1	-1	-1
-1	-1	-1	-1

$V_2$

0	-1	-2	-2
-1	-2	-2	-2
-2	-2	-2	-2
-2	-2	-2	-2

$V_3$

0	-1	-2	-3
-1	-2	-3	-3
-2	-3	-3	-3
-3	-3	-3	-3

$V_4$

0	-1	-2	-3
-1	-2	-3	-4
-2	-3	-4	-4
-3	-4	-4	-4

$V_5$

0	-1	-2	-3
-1	-2	-3	-4
-2	-3	-4	-5
-3	-4	-5	-5

$V_6$

0	-1	-2	-3
-1	-2	-3	-4
-2	-3	-4	-5
-3	-4	-5	-6

$V_7$

# Synchronous dynamic programming algorithms

Problem	Bellman Equation	Algorithm
Prediction	Bellman expectation equation	Iterative policy evaluation
Control	Bellman expectation equation +Greedy policy improvement	Policy iteration
Control	Bellman optimality equation	Value iteration

- > Algorithms are based on state-value function  $V_\pi(s)$  or  $V^*(s)$
- > Complexity  $O(mn^2)$  per iteration, for  $m$  actions and  $n$  states
- > Could also apply to action-value function  $Q_\pi(s, a)$  or  $Q^*(s, a)$
- > Complexity  $O(m^2n^2)$  per iteration

## Asynchronous dynamic programming

- > Policy/value iteration require exhaustive sweeps of the entire states.
- > Asynchronous DP backs up states individually, in any order
  - Sample a state at random and apply the appropriate backup
- > Can significantly reduce computation
- > Guaranteed to converge if **all states continue to be selected**

## Asynchronous dynamic programming



- > Three simple ideas for asynchronous dynamic programming:
  - In-place dynamic programming:  
only stores one copy of value function
  - Prioritized sweeping:  
use magnitude of Bellman error to guide state selection  
backup the state with the largest remaining Bellman error  
(requires knowledge of reverse dynamics)
  - Real-time dynamic programming:  
focus on states that are relevant to agent  
(use agent's experience to guide the selection of states)

## Full-width backups

- > Standard DP uses full-width backups
- > For each backup (sync or async)
  - every successor state and action is considered
  - using true model of transitions and reward function
- > DP is effective for medium-sized problems (millions of states)
- > For large problems, DP suffers from curse of dimensionality
  - even one full backup can be too expensive

## Sample backups

- > Sample backups: using sample rewards and sample transitions
  - Instead of reward function and transition dynamics
  - Advantages:
    - model-free: no prior knowledge of MDP required
    - breaks the curse of dimensionality through sampling
    - cost of backup is constant, independent of the state dimension



## Model-free prediction and control


- > Estimate or optimize the value function of an unknown MDP
- > MC and TD
  - learn directly from episodes of experience
  - no knowledge of MDP transitions / rewards
- > Monte-Carlo
  - learn from complete episodes: no bootstrapping
  - can only be applied to episodic MDPs (all episodes must terminate)
- > Temporal-difference
  - learn from incomplete episodes by bootstrapping
  - update a guess towards a guess

## Monte-Carlo policy evaluation

- > Goal: learn  $V_\pi$  from episodes of experience under policy  $\pi$

$$S_1, A_1, R_2, \dots, S_k \sim \pi$$

- > Monte-Carlo policy evaluation uses empirical mean return instead of expected return



## Monte-Carlo policy evaluation

- > To evaluate state  $s$
- > The first/every time-step  $t$  that state  $s$  is visited in an episode,
  - increment counter  $N(s) \leftarrow N(s) + 1$
  - increment total return  $S(s) \leftarrow S(s) + G_t$
  - value is estimated by mean return  $V(s) = S(s)/N(s)$
  - $V(s) \rightarrow V_\pi(s)$  as  $N(s) \rightarrow \infty$
- > Incremental updates
  - $V(s) \leftarrow V(s) + \frac{1}{N(s)}(G_t - V(s))$
  - To forget old episodes,  
$$V(s) \leftarrow V(s) + \alpha(G_t - V(s))$$

## TD learning


- > Incremental every-visit Monte-Carlo
  - update value toward actual return  $G_t$
  - $V(s) \leftarrow V(s) + \alpha(G_t - V(s))$
- > Simplest temporal-difference learning algorithm: TD(0)
  - update value towards estimated return  $r + \gamma V(s')$
  - $V(s) \leftarrow V(s) + \alpha(r + \gamma V(s') - V(s))$
  - TD error:  $r + \gamma V(s') - V(s)$

# TD learning

- > TD can learn before knowing the final outcome
  - can learn online after every step
  - MC must wait until end of episode before return is known
- > TD can learn without the final outcome
  - can learn from incomplete sequences
  - can work in continuing (non-terminating) environments
  - while MC can't
- > Disadvantage
  - Return  $G_t$  is unbiased estimate of  $V(s)$ , but TD target is biased
  - MC has high variance, zero bias -> good convergence
  - TD has low variance, some bias -> sensitive to initial value


# Unified view of RL

- > Bootstrapping:  
update involves an estimate
  - MC does not bootstrap
  - DP bootstrap
  - TD bootstrap
  
- > Sampling:  
update samples an expectation
  - MC samples
  - DP does not sample
  - TD samples



# TD( $\lambda$ )


- > n-step prediction: Let TD target look n steps into the future



- > Forward-view TD( $\lambda$ )  
$$V(s) \leftarrow V(s) + \alpha(G_t^\lambda - V(s))$$
- > Backward-view TD( $\lambda$ )

# On and off-policy learning

- > On-policy learning
  - learn about policy  $\pi$  from experience sampled from  $\pi$
- > Off-policy learning
  - learn about policy  $\pi$  from experience sampled from other policies



## On policy learning

- > Greedy policy improvement over  $Q(s, a)$  is model-free

$$\pi'(s) = \operatorname{argmax}_a Q(s, a)$$

- > Monte-Carlo policy iteration (Monte-Carlo policy eval. +  $\epsilon$ -greedy)
- > TD learning on action-value function: SARSA (Sarsa +  $\epsilon$ -greedy)

## Off policy learning

- > Evaluate target policy  $\pi(a|s)$  to compute  $V_\pi(s)$  or  $Q_\pi(s, a)$
- > Importance sampling: estimate the expectation of a different distribution

$$\begin{aligned}\mathbb{E}_{X \sim P}[f(X)] &= \sum P(X)f(X) \\ &= \sum Q(X) \frac{P(X)}{Q(X)} f(X) \\ &= \mathbb{E}_{X \sim Q} \left[ \frac{P(X)}{Q(X)} f(X) \right]\end{aligned}$$

- > Multiply importance sampling corrections along whole episode
  - can dramatically increase variance

$$G_t^{\pi/\mu} = \frac{\pi(A_t|S_t)}{\mu(A_t|S_t)} \frac{\pi(A_{t+1}|S_{t+1})}{\mu(A_{t+1}|S_{t+1})} \cdots \frac{\pi(A_T|S_T)}{\mu(A_T|S_T)} G_t$$

$$V(S_t) \leftarrow V(S_t) + \alpha \left( \textcolor{red}{G_t^{\pi/\mu}} - V(S_t) \right)$$

## Off policy learning

- > Q-learning (off-policy TD learning)
  - No importance sampling is required
  - Next action is chosen using behavior policy
  - but we consider alternative successor action  $a^+$ , and update  $Q(s, a)$  towards value of alternative action

$$Q(s, a) \leftarrow Q(s, a) + \alpha(r + \gamma Q(s', a^+) - Q(s, a))$$

- We allow both behavior policy and target policy to improve

$$Q(s, a) \leftarrow Q(s, a) + \alpha(r + \gamma \max_a Q(s', a') - Q(s, a))$$

- Q-learning control converges to the optimal action-value function

# Model-free control summary

	<i>Full Backup (DP)</i>	<i>Sample Backup (TD)</i>
Bellman Expectation Equation for $v_\pi(s)$	$v_\pi(s) \leftarrow s$ <p style="text-align: center;">Iterative Policy Evaluation</p>	$v_\pi(s') \leftarrow s'$ <p style="text-align: center;">TD Learning</p>
Bellman Expectation Equation for $q_\pi(s, a)$	$q_\pi(s, a) \leftarrow s, a$ <p style="text-align: center;">Q-Policy Iteration</p>	$q_\pi(s', a') \leftarrow a'$ <p style="text-align: center;">Sarsa</p>
Bellman Optimality Equation for $q_*(s, a)$	$q_*(s, a) \leftarrow s, a$ <p style="text-align: center;">Q-Value Iteration</p>	$q_*(s', a') \leftarrow a'$ <p style="text-align: center;">Q-Learning</p>

## Model-free control summary

<i>Full Backup (DP)</i>	<i>Sample Backup (TD)</i>
Iterative Policy Evaluation	TD Learning
$V(s) \leftarrow \mathbb{E}[R + \gamma V(S') \mid s]$	$V(S) \xleftarrow{\alpha} R + \gamma V(S')$
Q-Policy Iteration	Sarsa
$Q(s, a) \leftarrow \mathbb{E}[R + \gamma Q(S', A') \mid s, a]$	$Q(S, A) \xleftarrow{\alpha} R + \gamma Q(S', A')$
Q-Value Iteration	Q-Learning
$Q(s, a) \leftarrow \mathbb{E} \left[ R + \gamma \max_{a' \in \mathcal{A}} Q(S', a') \mid s, a \right]$	$Q(S, A) \xleftarrow{\alpha} R + \gamma \max_{a' \in \mathcal{A}} Q(S', a')$

where  $x \xleftarrow{\alpha} y \equiv x \leftarrow x + \alpha(y - x)$