

ECE7121 Learning-based control – 2025 Fall

# Policy gradient





INHA UNIVERSITY

# Overview

- > DQN
- > REINFORCE
- > Off-policy policy gradient

# Deep Q-learning

## > Deep Q-Network (DQN)



## Deep Q-learning

- > Epsilon greedy policy gives  $a$  and observe  $(s, a, s', r)$
- > Q-learning

$$Q(s, a) \leftarrow Q(s, a) + \alpha(r + \gamma \max_a Q(s', a') - Q(s, a))$$

- > DQN, Q function is represented by neural network

- $Q \sim Q_\phi$
- sample from the buffer and compute the target


$$y_i = r(s_i, a_i) + \gamma \max_a Q_{\phi^-}(s'_i, a'_i)$$

- update


$$\phi \leftarrow \operatorname{argmin}_\phi \sum_i \|Q_\phi(s_i, a_i) - y_i\|^2$$

# Deep Q-learning

## Classic Q-learning



## Deep Q-learning



## Deep Q-learning

- > Convergence guarantee under some assumptions
  - data collection policy has good coverage
  - realizability:  $Q_\phi$  can represent  $Q$
  - $\phi$  has some good property
- > Online update
  - $\phi \leftarrow \operatorname{argmin}_\phi \sum_i \|Q_\phi(s_i, a_i) - y_i\|^2$
  - $\phi \leftarrow \phi - \eta \sum_i (Q_\phi(s_i, a_i) - y_i) \nabla_\phi Q_\phi(s_i, a_i)$
- > DQN paper ideas (DeepMind 2013, later 2015 nature)
  - Replay buffer: Q-learning is off-policy! save all transition data  $(s, a, s', r)$
  - target network: use two Q-networks, periodically update parameters

# Deep Q-learning

- > Practical tips (general RL tips)
  - it can take a while to converge (be patient)
  - high exploration initially, then reduce
  - test on known / easy tasks first
  - use / add more advanced features
  - use a well tested library (e.g., stable-baselines3)

## DQN extensions

- > DDQN
- > Prioritized experience replay
- > Dueling architecture
- > **Q-learning with continuous actions -> DDPG**

# Q-learning with continuous actions



- > How to solve max or argmax with continuous actions?

$$\mathbb{E}_{(s,a,s',r)}[(Q_\phi(s, a) - r(s, a) - \gamma \max_{a'} Q_\phi^-(s', a'))^2]$$

- > Option 1: discretization
  - not scalable
- > Option 2: optimization
  - gradient-based optimization (GD, SGD) on  $Q_\phi(s', a')$
  - stochastic optimization
    - Cross-entropy method (CEM): simple iterative stochastic optimization
    - CMA-ES: more complex iterative stochastic optimization
  - slow and stuck in the local minimum

# Stochastic optimization

> CEM / CMA-ES



## Q-learning with continuous actions


- > How to solve max or argmax with continuous actions?

$$\mathbb{E}_{(s,a,s',r)}[(Q_\phi(s, a) - r(s, a) - \gamma \max_{a'} Q_{\phi^-}(s', a'))^2]$$

- > Option 3: structured Q functions that are easy to optimize
  - example: normalized advantage functions (NAF)
  - problem: less expressive power

$$Q_\phi(\mathbf{s}, \mathbf{a}) = -\frac{1}{2}(\mathbf{a} - \mu_\phi(\mathbf{s}))^T P_\phi(\mathbf{s})(\mathbf{a} - \mu_\phi(\mathbf{s})) + V_\phi(\mathbf{s})$$

$$\arg \max_{\mathbf{a}} Q_\phi(\mathbf{s}, \mathbf{a}) = \mu_\phi(\mathbf{s}) \quad \max_{\mathbf{a}} Q_\phi(\mathbf{s}, \mathbf{a}) = V_\phi(\mathbf{s})$$



# Q-learning with continuous actions

- > How to solve max or argmax with continuous actions?


$$\mathbb{E}_{(s,a,s',r)} \left[ \left( Q_\phi(s, a) - r(s, a) - \gamma \max_{a'} Q_\phi(s', a') \right)^2 \right]$$

- > **Option 4: learn an approximate optimizer**

- DDPG (deep deterministic policy gradient)
- Main idea: train another DNN  $\mu_\theta(s)$  such that  
 $\mu_\theta(s) \approx \arg \max_a Q_\phi(s, a)$
- $\mu_\theta(s)$  is essentially a policy
- How? gradient ascent on  $Q_\phi(s', a')$ , using a chain rule  $\frac{dQ_\phi}{d\phi} = \frac{dQ_\phi}{da} \frac{da}{d\theta}$
- Policy gradient: can we directly optimize  $\mu_\theta(s)$ ?

## Summary of Q-learning


- > Key idea:




- > Q-learning is off-policy and model-free
- > Q-learning doesn't explicitly learn a policy
- > Q-learning could be hard with continuous actions
- > Online Q-learning with neural network approximations: DQN

# Why policy gradient?

- > Key philosophy: why learn value function when all we need is a policy?
  - In some cases, policy might be easier to learn!
  - ex) In LQR, the optimal policy is linear while the optimal value is quadratic
- > DQN:



- > Policy gradient: directly optimize the policy  $\pi_\phi$  via  $\nabla_\phi J(\phi)$




# Policy functions

## > Advantages over value based methods

- effective in high-dimensional or continuous action spaces
- can learn stochastic policies


deterministic continuous policy



$$a = \pi_\theta(s)$$

e.g. outputs a steering angle directly

stochastic continuous policy



$$a \sim \mathcal{N}(\mu_\theta(s), \sigma_\theta^2(s))$$

FA for stochastic multimodal continuous policies is an active area of research

(stochastic) policy over discrete actions



go left  
go right  
press brake

Outputs a distribution over a discrete set of actions

## Policy gradient theorem

- > Policy gradient: directly optimize the policy  $\pi_\phi$  via  $\nabla_\phi J(\phi)$

- $p_\phi(s_1, a_1, \dots, s_T, a_T) = p(s_1) \prod_{t=1}^T \pi_\phi(a_t | s_t) p(s_{t+1} | s_t, a_t)$

$$\phi^* = \underset{\underbrace{J(\phi)}}{\operatorname{argmax}_\phi \mathbb{E}_{\tau \sim p_\phi(\tau)} [\sum_t r(s_t, a_t)]}$$

- > Given some policy  $\pi_\phi$ , how to evaluate its  $J(\phi)$
- > Simple idea: Monte-Carlo sampling
  - sample N trajectories using  $\pi_\phi$
  - Monte-Carlo estimation of  $J(\phi)$ ,  $J(\phi) = \frac{1}{N} \sum_i \sum_t r(s_{i,t}, a_{i,t}, t)$
- > Can we estimate  $\nabla_\phi J(\phi)$ ?

## Policy gradient theorem

- > How to compute  $\nabla_{\phi}J(\phi)$ ?
  - $\nabla_{\phi}J(\phi)$  is hard to compute:  $\phi$  influences both the action selections and the distribution of states in which those selections are made
- > Policy gradient theorem:

$$\nabla_{\phi}J(\phi) = E_{\tau \sim p_{\phi}(\tau)}[R(\tau) \sum_t \nabla_{\phi} \log \pi_{\phi}(a_t | s_t)]$$

$$R(\tau) = \sum_t r(s_t, a_t)$$

# Policy gradient theorem

## > Derivation

- $J(\phi) = \mathbb{E}_{\tau \sim p_\phi(\tau)} [\sum_t r(s_t, a_t)] = \int p_\phi(\tau) R(\tau) d\tau$
- $\nabla_\phi J(\phi) = \int \nabla_\phi p_\phi(\tau) R(\tau) d\tau = \int p_\phi(\tau) \nabla_\phi \log p_\phi(\tau) R(\tau) d\tau$   
 $= \mathbb{E}_{\tau \sim p_\phi(\tau)} [\nabla_\phi \log p_\phi(\tau) R(\tau)]$   
*since  $p_\phi(\tau) \nabla_\phi \log p_\phi(\tau) = p_\phi(\tau) \frac{\nabla_\phi p_\phi(\tau)}{p_\phi(\tau)} = \nabla_\phi p_\phi(\tau)$*
- $\nabla_\phi \log p_\phi(\tau) = \nabla_\phi \log p(s_1) \prod_{t=1}^T \pi_\phi(a_t | s_t) p(s_{t+1} | s_t, a_t)$   
 $= \nabla_\phi \log p(s_1) + \sum_{t=1}^T \log \pi_\phi(a_t | s_t) + \log p(s_{t+1} | s_t, a_t)$   
 $= \nabla_\phi \sum_{t=1}^T \log \pi_\phi(a_t | s_t)$   
 $= \sum_{t=1}^T \nabla_\phi \log \pi_\phi(a_t | s_t)$
- $\nabla_\phi J(\phi) = \mathbb{E}_{\tau \sim p_\phi(\tau)} [\sum_{t=1}^T \nabla_\phi \log \pi_\phi(a_t | s_t) R(\tau)]$

# REINFORCE


## > REINFORCE = Monte-Carlo PG

- estimate the expectation by Monte-Carlo sampling

$$\begin{aligned}\nabla_{\phi} J(\phi) &= \mathbb{E}_{\tau \sim p_{\phi}(\tau)} \left[ \sum_{t=1}^T \nabla_{\phi} \log \pi_{\phi}(a_t | s_t) R(\tau) \right] \\ &\approx \frac{1}{N} \sum_i R(\tau^i) (\sum_t \nabla_{\phi} \log \pi_{\phi}(a_{i,t} | s_{i,t}))\end{aligned}$$

- Algorithm (Sutton et al., NeurIPS 1998)

1. sample  $\{\tau^i\}$  from  $\pi_{\phi}(a_t | s_t)$  (run the policy)
2. estimate  $\nabla_{\phi} J(\phi)$
3.  $\phi \leftarrow \phi + \alpha \nabla_{\phi} J(\phi)$  and go to 1



# Understanding REINFORCE

- > In imitation learning



maximum likelihood:  $\nabla_\phi J_{ML}(\theta) \approx \frac{1}{N} \sum_{i=1}^N (\sum_{t=1}^T \nabla_\phi \log \pi_\phi(a_{i,t} | s_{i,t}))$

- maximizing log likelihood of the policy  $\pi_\phi$  on expert data

- > Policy gradient:

- imitation gradient, but weighted by reward

$$\nabla_\phi J(\phi) \approx \frac{1}{N} \sum_{i=1}^N R(\tau^i) (\sum_{t=1}^T \nabla_\phi \log \pi_\phi(a_{i,t} | s_{i,t}))$$

# Understanding REINFORCE

> Policy gradient:

$$\nabla_{\phi} J(\phi) \approx \frac{1}{N} \sum_{i=1}^N R(\tau^i) (\sum_{t=1}^T \nabla_{\phi} \log \pi_{\phi}(a_{i,t} | s_{i,t}))$$

- with out expert data: signal is from  $R(\tau^i)$
- Good stuff (bigger  $R(\tau^i)$ ) is made more likely
- Bad stuff is made less likely
- formalization of trial-and-error
  
- It is on-policy ( $\mathbb{E}_{\tau \sim p_{\phi}(\tau)}$ ) and model-free (no use of  $p(s_{t+1} | s_t, a_t)$ )
- Markov property is not used  $\Rightarrow$  it can be used in POMDP

# Understanding REINFORCE

- > Policy gradient:

$$\nabla_{\phi} J(\phi) \approx \frac{1}{N} \sum_{i=1}^N R(\tau^i) (\sum_{t=1}^T \nabla_{\phi} \log \pi_{\phi}(a_{i,t} | s_{i,t}))$$

- > Weakness
  - It has high variance because it relies on full-trajectory
  - There might be good samples, but have  $R(\tau^i) = 0$
  - It requires waiting until the end of the episodes

## Reducing variance

- >  $\nabla_{\phi} J(\phi) \approx \frac{1}{N} \sum_{i=1}^N (\sum_t \nabla_{\phi} \log \pi_{\phi}(a_{i,t} | s_{i,t}) \sum_t r(s_{i,t}, a_{i,t}))$
- > Causality: policy at time  $t'$  cannot affect reward  $t < t'$

$$\nabla_{\phi} J(\phi) \approx \frac{1}{N} \sum_{i=1}^N \sum_t \nabla_{\phi} \log \pi_{\phi}(a_{i,t} | s_{i,t}) \underbrace{(\sum_{t'=t} r(s_{i,t}, a_{i,t}))}_{\begin{array}{c} \text{Rewards to go} \\ Q^{\pi}(s, a) \end{array}}$$

- > Another form of PG theorem, using  $Q^{\pi}$  to estimate return

$$\nabla_{\phi} J(\phi) = \mathbb{E}_{s \sim p_{\phi}, a \sim \pi_{\phi}} [\sum_{t=1}^T \nabla_{\phi} \log \pi_{\phi}(a_t | s_t) Q^{\pi}(s, a)]$$

## Reducing variance

- > Add a bias term does not affect gradient

$$\begin{aligned} & \mathbb{E}_{s \sim p_\phi, a \sim \pi_\phi} \left[ \sum_{t=1}^T \nabla_\phi \log \pi_\phi(a_t | s_t) b(s) \right] \\ &= \iint p_\phi(s) \pi_\phi(a|s) \frac{\nabla_\phi \pi_\phi(a|s)}{\pi_\phi(a|s)} b(s) dads \\ &= \iint p_\phi(s) \nabla_\phi \pi_\phi(a|s) b(s) dads \\ &= \int p_\phi(s) b(s) \nabla_\phi \int \pi_\phi(a|s) da ds \\ &= \int p_\phi(s) b(s) \nabla_\phi 1 ds = 0 \end{aligned}$$

- > Therefore

$$\nabla_\phi J(\phi) = \mathbb{E}_{s \sim p_\phi, a \sim \pi_\phi} \left[ \sum_{t=1}^T \nabla_\phi \log \pi_\phi(a_t | s_t) (Q^\pi(s, a) - b(s)) \right]$$

- average reward or state value function can be used for  $b(s)$
- optimal bias is the expected reward weighted by gradient magnitudes




## Off-policy PG

- > Neural networks change only a little bit with each gradient step
  - since we have a tons of parameters
- > On-policy learning can be extremely inefficient!
  - cannot reuse old data
- > Algorithm
  1. sample  $\{\tau^i\}$  from  $\pi_\phi(a_t|s_t)$  (run the policy)
  2. estimate  $\nabla_\phi J(\phi)$
  3.  $\phi \leftarrow \phi + \alpha \nabla_\phi J(\phi)$  and go to 1       $\leftarrow$  we change  $\phi$  here

# Importance sampling

> How can we use the data from other policies?




$$\begin{aligned} - \mathbb{E}_{x \sim p(x)}[f(x)] &= \int p(x)f(x)dx \\ &= \int \frac{q(x)}{q(x)}p(x)f(x)dx \\ &= \int q(x)\frac{p(x)}{q(x)}f(x)dx \\ &= \mathbb{E}_{x \sim q(x)}\left[\frac{p(x)}{q(x)}f(x)\right] \end{aligned}$$



# Importance sampling

> How can we use the data from other policies?

- $\mathbb{E}_{x \sim p(x)}[f(x)] = \mathbb{E}_{x \sim q(x)} \left[ \frac{p(x)}{q(x)} f(x) \right]$



## Off-policy PG

> How can we use the data from other policies?

- $\mathbb{E}_{x \sim p(x)}[f(x)] = \mathbb{E}_{x \sim q(x)} \left[ \frac{p(x)}{q(x)} f(x) \right]$

> Likewise, we have samples from different policy  $p_{\phi'}(\tau)$

- $J(\phi') = \mathbb{E}_{\tau \sim p_{\phi'}(\tau)}[R(\tau)] = \mathbb{E}_{\tau \sim p_{\phi}(\tau)} \left[ \frac{p_{\phi'}(\tau)}{p_{\phi}(\tau)} R(\tau) \right]$

- $\nabla_{\phi'} J(\phi') = \mathbb{E}_{\tau \sim p_{\phi}(\tau)} \left[ \nabla_{\phi'} \frac{p_{\phi'}(\tau)}{p_{\phi}(\tau)} R(\tau) \right]$

$$= \mathbb{E}_{\tau \sim p_{\phi}(\tau)} \left[ \frac{p_{\phi'}(\tau)}{p_{\phi}(\tau)} \nabla_{\phi'} \log p_{\phi'}(\tau) R(\tau) \right]$$

$$= \mathbb{E}_{\tau \sim p_{\phi}(\tau)} \left[ \frac{p_{\phi'}(\tau)}{p_{\phi}(\tau)} (\sum_t \nabla_{\phi'} \log \pi_{\phi'}(a_t | s_t)) (\sum_t r(s_t, a_t)) \right]$$

## Off-policy PG

- $\nabla_{\phi'} J(\phi') = \mathbb{E}_{\tau \sim p_\phi(\tau)} \left[ \frac{p_{\phi'}(\tau)}{p_\phi(\tau)} (\sum_t \nabla_{\phi'} \log \pi_{\phi'}(a_t|s_t)) (\sum_t r(s_t, a_t)) \right]$
- $\frac{p_{\phi'}}{p_\phi} = \frac{p(s_1) \prod_t \pi_{\phi'}(a_t|s_t) p(s_{t+1}|s_t, a_t)}{p(s_1) \prod_t \pi_\phi(a_t|s_t) p(s_{t+1}|s_t, a_t)} = \frac{\prod_t \pi_{\phi'}(a_t|s_t)}{\prod_t \pi_\phi(a_t|s_t)}$
- $\nabla_{\phi'} J(\phi') = \mathbb{E}_{\tau \sim p_\phi(\tau)} \left[ \underbrace{\left( \frac{\prod_t \pi_{\phi'}(a_t|s_t)}{\prod_t \pi_\phi(a_t|s_t)} \right)}_{\text{Causality}} (\sum_t \nabla_{\phi'} \log \pi_{\phi'}(a_t|s_t)) (\sum_t r(s_t, a_t)) \right]$
- Causality  
 $= \mathbb{E}_{\tau \sim p_\phi(\tau)} \left[ \underbrace{\left( \frac{\prod_t \pi_{\phi'}(a_t|s_t)}{\prod_t \pi_\phi(a_t|s_t)} \right)}_{\text{Causality}} (\sum_t \nabla_{\phi'} \log \pi_{\phi'}(a_t|s_t)) (\sum_{t'=t} r(s_t, a_t)) \right]$ 

↓
- Problem: orange is exponential in horizon, can become very small or very large. Also hard to measure

## Off-policy PG

### > Reducing variance

- $\mathbb{E}_{\tau \sim p_{\phi}(\tau)} \left[ \underbrace{\left( \frac{\prod_t \pi_{\phi'}(a_t | s_t)}{\prod_t \pi_{\phi}(a_t | s_t)} \right) (\sum_t \nabla_{\phi'} \log \pi_{\phi'}(a_t | s_t)) (\sum_{t'=t} r(s_t, a_t) - b)}_{\text{redacted}} \right]$

### > Instead of using trajectories, consider the expectation over timesteps

- $\nabla_{\phi} J(\phi') \approx \frac{1}{N} \sum_i \sum_t \left( \frac{\pi_{\phi'}(a_{i,t} | s_{i,t})}{\pi_{\phi}(a_{i,t} | s_{i,t})} \right) (\nabla_{\phi'} \log \pi_{\phi'}(a_{i,t} | s_{i,t})) ((\sum_{t'=t} r(s_{i,t}, a_{i,t}) - b)$

- this would be the common final form
- we can take multiple gradients steps on the same batch

### > Algorithm

1. sample  $\{\tau^i\}$  from  $\pi_{\phi}(a_t | s_t)$  (run the policy)
2. estimate  $\nabla_{\phi} J(\phi)$
3.  $\phi \leftarrow \phi + \alpha \nabla_{\phi} J(\phi)$  and go to 1

can take multiple gradient steps

## Off-policy PG

- > What if our policy changes a lot before sampling new data?
  - Data no longer reflects states that policy will visit
  - gradient estimate less accurate
- > Can we constrain the policy not stray too far during gradient update?
  - One common choice:  $\mathbb{E}_{s \sim \pi_\phi} [D_{KL} (\pi_{\phi'}(\cdot | s) \| \pi_\phi(\cdot | s))] \leq \delta$
  - We will learn in the future

## Review

- > Since the gradient has high variance,
  - consider using much larger batches
  - tuning the learning rates is challenging
  - Adaptive step size rules like Adam can be used
- > Policy gradient is on-policy
  - off-policy variant using importance sampling
  - Some approximations are made by ignoring parts of the state
  - We can apply multiple gradients updates
- > Intuition
  - do more high reward stuff, less low reward stuff
  - Gradient will be very noisy, best with large batch sizes and dense rewards