

ECE7121 Learning-based control – 2025 Fall

Actor-critic



INHA UNIVERSITY

Overview

- > Advantage
- > Actor-critic structure
- > PPO
- > Off-policy AC

Recap

> Policy gradients

$$\nabla_{\phi} J(\phi) \approx \frac{1}{N} \sum_{i=1}^N \sum_t \nabla_{\phi} \underbrace{\log \pi_{\phi}(a_{i,t} | s_{i,t})}_{\text{policy likelihood}} \underbrace{(\sum_{t'=t} r(s_{i,t}, a_{i,t}) - b)}_{\text{rewards to go}} - \underbrace{b}_{\text{baseline}}$$

samples from policy

- do more of the above average stuff, less of the below average stuff



> Importance weights for off-policy PG

$$\nabla_{\phi'} J(\phi') \approx \frac{1}{N} \sum_i \sum_t \underbrace{\left(\frac{\pi_{\phi'}(a_{i,t} | s_{i,t})}{\pi_{\phi}(a_{i,t} | s_{i,t})} \right)}_{\text{importance weight}} (\nabla_{\phi'} \log \pi_{\phi'}(a_{i,t} | s_{i,t})) ((\sum_{t'=t} r(s_{i,t}, a_{i,t}) - b))$$

- only suitable when policies are very similar

Advantage

- > Value function $V^\pi(s)$
 - future expected rewards starting at s and following π
- > State-action value function $Q^\pi(s, a)$
 - future expected rewards starting at s , taking a , then following π



$$V^\pi(s) = \mathbb{E}_{a \sim \pi(\cdot|s)}[Q^\pi(s, a)]$$


- > Advantage $A^\pi(s, a)$
 - how much better it is to take a than to follow π at state s
 - $A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s)$

Advantage

- > $\nabla_{\phi} J(\phi) \approx \frac{1}{N} \sum_{i=1}^N \sum_t \nabla_{\phi} \log \pi_{\phi}(a_{i,t} | s_{i,t}) (\sum_{t'=t} r(s_{i,t}, a_{i,t}) - b)$
 $\rightarrow \frac{1}{N} \sum_{i=1}^N \sum_t \nabla_{\phi} \log \pi_{\phi}(a_{i,t} | s_{i,t}) (Q(s_{i,t}, a_{i,t}) - b)$
 $\rightarrow \frac{1}{N} \sum_{i=1}^N \sum_t \nabla_{\phi} \log \pi_{\phi}(a_{i,t} | s_{i,t}) (Q(s_{i,t}, a_{i,t}) - V^{\pi}(s_t))$
 $= \frac{1}{N} \sum_{i=1}^N \sum_t \nabla_{\phi} \log \pi_{\phi}(a_{i,t} | s_{i,t}) A^{\pi}(s_{i,t}, a_{i,t})$
- > Better estimates of A^{π} lead to less noisy gradients

Advantage

- > Estimating expected return
- > Should we fit V^π , Q^π or A^π ?
- >
$$\begin{aligned} A^\pi(s_t, a_t) &= Q^\pi(s_t, a_t) - V^\pi(s_t) \\ &= \sum_{t'=t}^T \mathbb{E}_\pi[r(s_{t'}, a_{t'})|s_t, a_t] - V^\pi(s_t) \\ &= r(s_t, a_t) + \sum_{t'=t+1}^T \mathbb{E}_\pi[r(s_{t'}, a_{t'})|s_t, a_t] - V^\pi(s_t) \\ &= r(s_t, a_t) + \mathbb{E}_{s_{t+1} \sim p(\cdot|s_t, a_t)}[V^\pi(s_{t+1})] - V^\pi(s_t) \\ &\approx r(s_t, a_t) + V^\pi(s_{t+1}) - V^\pi(s_t) \end{aligned}$$
- > Let's just fit V^π



Estimating the value function

> In policy evaluation,


- Monte-Carlo: $V(s) \leftarrow V(s) + \alpha(G_t - V(s))$
- TD(0): $V(s) \leftarrow V(s) + \alpha(r + \gamma V(s') - V(s))$
- n-step TD: $V(s) \leftarrow V(s) + \alpha(G_t^{(n)} - V(s))$
 - $G_t^{(n)} = r_t + \gamma r_{t+1} + \dots + \gamma^{n-1} r_{t+n-1} + \gamma^n V(s_{t+n})$

> Likewise, fitting can be done in a similar way

- Monte-Carlo: $\hat{V}_\theta^\pi(s_i) \leftarrow \sum_{t'=t}^T r(s_{i,t'}, a_{i,t'})$
- TD(0): $\hat{V}_\theta^\pi(s_i) \leftarrow r(s_{i,t'}, a_{i,t'}) + \gamma \hat{V}_\theta^\pi(s_{i,t+1})$
- n-step TD: $\hat{V}_\theta^\pi(s_i) \leftarrow \sum_{t'=t}^{t+n-1} \gamma^{t'-t} r(s_{i,t'}, a_{i,t'}) + \gamma^n \hat{V}_\theta^\pi(s_{i,t+n})$

Estimating the value function

- > Monte-Carlo: supervise with roll-out's summed rewards
- > Bootstrap: supervise using reward and current value estimate
- > N-step returns: lies between MC and Bootstrapped
 - less variance than MC
 - lower bias than 1-step bootstrap
 - often works the best in practice





A full actor-critic algorithm

> Actor-critic: learn to estimate what is good vs. bad, then do more of the good stuff

- get a better PG by using NN to estimate value function



> Algorithm

1. sample a batch of data $\{\tau^i\}$ from $\pi_\phi(a_t|s_t)$ (run the policy)
2. Fit $\hat{V}_\theta^{\pi_\phi}$ to summed rewards in data (n-step TD)
3. Evaluate $\hat{A}^{\pi_\phi}(s_{i,t}, a_{i,t}) = r(s_{i,t}, a_{i,t}) + \gamma \hat{V}_\theta^{\pi_\phi}(s_{i,t+1}) - \hat{V}_\theta^{\pi_\phi}(s_{i,t})$
4. Evaluate $\nabla_\phi J(\phi) \approx \sum_{i=1}^N \sum_t \nabla_\phi \log \pi_\phi(a_{i,t}|s_{i,t}) \hat{A}^{\pi_\phi}(s_{i,t}, a_{i,t})$
5. $\phi \leftarrow \phi + \alpha \nabla_\phi J(\phi)$ and go to 1




Architecture design

- > Two network design



- Simple and stable, but no shared features between actor & critic


- > Shared network design



- computational efficiency, shared representation learning, good empirical performance

Online actor-critic in practice

- > Online actor-critic works best with multiple parallel environments
 - consecutive observations are temporally correlated, violating i.i.d.
 - diversity of data allows the agent to generalize better and avoid overfitting
 - improve the signal-to-noise ratio of policy updates
- > Updates can be done either synchronously or asynchronously



Off-policy AC

> Version 1. Use multiple gradient steps using IS

1. sample a batch of data $\{\tau^i\}$ from $\pi_\phi(a_t|s_t)$ (run the policy)

2. fit $\hat{V}_\theta^{\pi_\phi}$ to summed rewards in data (n-step TD)

3. evaluate $\hat{A}^{\pi_\phi}(s_{i,t}, a_{i,t}) = r(s_{i,t}, a_{i,t}) + \gamma \hat{V}_\theta^{\pi_\phi}(s_{i,t+1}) - \hat{V}_\theta^{\pi_\phi}(s_{i,t})$

4. evaluate $\nabla_{\phi'} J(\phi') \approx \sum_{i=1}^N \sum_t \left(\frac{\pi_{\phi'}(a_{i,t}|s_{i,t})}{\pi_\phi(a_{i,t}|s_{i,t})} \right) \nabla_{\phi'} \log \pi_{\phi'}(a_{i,t}|s_{i,t}) \hat{A}^{\pi_\phi}(s_{i,t}, a_{i,t})$

5. $\phi \leftarrow \phi + \alpha \nabla_{\phi} J(\phi)$

Advantages based on old policy
> will get out-dated

can take multiple gradient steps

> So far, we used one batch of policy data for multiple gradient steps

> Can we reuse data from previous batches?

- all the past trial-and-error data

Off-policy AC

> Version 1. Use multiple gradient steps using IS

1. sample a batch of data $\{\tau^i\}$ from $\pi_\phi(a_t|s_t)$ (run the policy)

2. fit $\hat{V}_\theta^{\pi_\phi}$ to summed rewards in data (n-step TD)

3. evaluate $\hat{A}^{\pi_\phi}(s_{i,t}, a_{i,t}) = r(s_{i,t}, a_{i,t}) + \gamma \hat{V}_\theta^{\pi_\phi}(s_{i,t+1}) - \hat{V}_\theta^{\pi_\phi}(s_{i,t})$

4. evaluate $\nabla_{\phi'} J(\phi') \approx \sum_{i=1}^N \sum_t \left(\frac{\pi_{\phi'}(a_{i,t}|s_{i,t})}{\pi_\phi(a_{i,t}|s_{i,t})} \right) \nabla_{\phi'} \log \pi_{\phi'}(a_{i,t}|s_{i,t}) \hat{A}^{\pi_\phi}(s_{i,t}, a_{i,t})$

5. $\phi \leftarrow \phi + \alpha \nabla_{\phi} J(\phi)$

Advantages based on old policy
> will get out-dated

can take multiple gradient steps

> Idea1: Use KL constraint on policy

- $\mathbb{E}_{s \sim \pi_\phi} [D_{KL} (\pi_{\phi'}(\cdot|s) \| \pi_\phi(\cdot|s))] \leq \delta$

- Trust region policy optimization (TRPO)

- very common. used in LLM preference optimization

Natural policy gradient

> Reformulate to a constrained optimization problem


- $\max_{\phi'} (\phi' - \phi)^\top \nabla_\phi J(\phi)$
 $s.t. \|\phi - \phi'\|^2 \leq \delta$

- $\max_{\phi'} (\phi' - \phi)^\top \nabla_\phi J(\phi)$
 $s.t. (\phi' - \phi)^\top F(\phi' - \phi) \leq \delta$


$$F = \mathbb{E}_{s \sim \pi_\phi} [\nabla_\phi \log \pi_\phi(a|s) \nabla_\phi \log \pi_\phi(a|s)^\top]$$

- Update becomes $\phi' = \phi + \alpha F^{-1} \nabla_\phi J(\phi)$

(a) ‘Vanilla’ policy gradients



(b) Natural policy gradients



Trust region policy optimization (TRPO)

> Constrained optimization problem

- $\max_{\phi'} (\phi' - \phi)^\top \nabla_\phi J(\phi)$

s.t. $\mathbb{E}_{s \sim \pi_\phi} [D_{KL}(\pi_{\phi'}(\cdot | s) \| \pi_\phi(\cdot | s))] \leq \delta$

$$D_{KL}(\pi_{\phi'}(\cdot | s) \| \pi_\phi(\cdot | s)) = \frac{1}{2} \Delta \phi^\top \underline{\nabla_{\phi'}^2 D_{KL}(\pi_{\phi'}(\cdot | s) \| \pi_\phi(\cdot | s)) \Delta \phi} F$$

- Update $\phi' = \phi + \eta F^{-1} \nabla_\phi J(\phi)$, $\eta = \sqrt{\frac{2\delta}{\nabla_\phi J^\top F^{-1} \nabla_\phi J}}$ (to meet the constraint)
- it is a constrained non-convex optimization problem
- complicated and hard/slow to solve
- PPO designs a surrogate objective function

Off-policy AC

> Version 1. Use multiple gradient steps using IS

1. sample a batch of data $\{\tau^i\}$ from $\pi_\phi(a_t|s_t)$ (run the policy)

2. fit $\hat{V}_\theta^{\pi_\phi}$ to summed rewards in data (n-step TD)

3. evaluate $\hat{A}^{\pi_\phi}(s_{i,t}, a_{i,t}) = r(s_{i,t}, a_{i,t}) + \gamma \hat{V}_\theta^{\pi_\phi}(s_{i,t+1}) - \hat{V}_\theta^{\pi_\phi}(s_{i,t})$

4. evaluate $\nabla_{\phi'} J(\phi') \approx \sum_{i=1}^N \sum_t \left(\frac{\pi_{\phi'}(a_{i,t}|s_{i,t})}{\pi_\phi(a_{i,t}|s_{i,t})} \right) \nabla_{\phi'} \log \pi_{\phi'}(a_{i,t}|s_{i,t}) \hat{A}^{\pi_\phi}(s_{i,t}, a_{i,t})$

5. $\phi \leftarrow \phi + \alpha \nabla_{\phi} J(\phi)$

Advantages based on old policy
> will get out-dated

can take multiple gradient steps

> Idea2: Can we bound the importance weights?

- Doesn't directly constrain policy, but removes incentives
- A key idea behind proximal policy optimization (PPO)

Proximal policy optimization (PPO)

$$> \nabla_{\phi'} J(\phi') \approx \sum_{i=1}^N \sum_t \left(\frac{\pi_{\phi'}(a_{i,t}|s_{i,t})}{\pi_{\phi}(a_{i,t}|s_{i,t})} \right) \nabla_{\phi'} \log \pi_{\phi'}(a_{i,t}|s_{i,t}) \hat{A}^{\pi_{\phi}}(s_{i,t}, a_{i,t})$$
$$J(\phi') \approx \sum_{t,i} \frac{\pi_{\phi'}(a_{i,t}|s_{i,t})}{\pi_{\phi}(a_{i,t}|s_{i,t})} \hat{A}^{\pi_{\phi}}(s_{i,t}, a_{i,t})$$
$$\underline{\hspace{10em}} [1] \qquad p_{\theta}(x) \nabla_{\theta} (\log p_{\theta}(x)) = \nabla_{\theta} p_{\theta}(x)$$

> Trick 1: clip the importance weights:

$$- J(\phi') \approx \sum_{t,i} \underline{\text{clip}\left(\frac{\pi_{\phi'}(a_{i,t}|s_{i,t})}{\pi_{\phi}(a_{i,t}|s_{i,t})}, 1 - \epsilon, 1 + \epsilon\right)} \hat{A}^{\pi_{\phi}}(s_{i,t}, a_{i,t})$$
$$\underline{\hspace{10em}} [2]$$

- policy no longer incentivized to deviate significantly

> Trick 2: take minimum w.r.t. original objective

- in rare event where clipping makes objective better
- $J(\phi') \approx \sum_{t,i} \min([1], [2])$

Proximal policy optimization (PPO)

- > $J(\phi') \approx \sum_{t,i} \min([1], [2])$
- > Trick 3: generalized advantage estimation (GAE)
 - Fit V^π with Monte-Carlo or bootstrapping
e.g. n-step TD: $\hat{V}_\theta^\pi(s_i) \leftarrow \sum_{t'=t}^{t+n-1} \gamma^{t'-t} r(s_{i,t'}, a_{i,t'}) + \gamma^n \hat{V}_\theta^\pi(s_{i,t+n})$
 - Then, use varying horizon to estimate advantage:
$$\hat{A}^{\pi_\phi}(s_{i,t}, a_{i,t}) = r(s_{i,t}, a_{i,t}) + \gamma \hat{V}_\theta^{\pi_\phi}(s_{i,t+1}) - \hat{V}_\theta^{\pi_\phi}(s_{i,t})$$
$$\rightarrow \hat{A}_n^{\pi_\phi}(s_{i,t}, a_{i,t}) = \sum_{t'=t}^{t+n} \gamma^{t'-t} r(s_{i,t'}, a_{i,t'}) + \gamma^n \hat{V}_\theta^{\pi_\phi}(s_{i,t+1}) - \hat{V}_\theta^{\pi_\phi}(s_{i,t})$$

GAE: $\hat{A}_{GAE}^{\pi_\phi}(s_{i,t}, a_{i,t}) = \sum_{n=1}^{\infty} \omega_n \hat{A}_n^{\pi_\phi}(s_{i,t}, a_{i,t})$
$$\omega_n \propto \lambda^{n-1}, \lambda \in (0,1)$$

Proximal policy optimization (PPO)


> Algorithm:

1. sample a batch of data $\{\tau^i\}$ from $\pi_\phi(a_t|s_t)$ (run the policy)
2. fit $\hat{V}_\theta^{\pi_\phi}$ to summed rewards in data (n-step TD)
- 3. evaluate $\hat{A}_{GAE}^{\pi_\phi}(s_{i,t}, a_{i,t})$**
- 4. compute the clipped surrogate objective gradient $\nabla_{\phi'} J(\phi')$**
5. update policy with M gradient steps $\phi \leftarrow \phi + \alpha \nabla_\phi J(\phi)$

Off-policy AC

> Version 2. Replay buffer


1. collect experience $\{s_i, a_i\}$ from π_ϕ and add to replay buffer
2. sample a batch $\{s_i, a_i, r_i, s_i'\}$ from buffer \mathcal{R}
3. update $\hat{V}_\theta^{\pi_\phi}$ using targets $r(s_i, a_i) + \gamma \hat{V}_\theta^\pi(s_i')$ for each s_i
4. evaluate $\hat{A}^{\pi_\phi}(s_i, a_i) = r(s_i, a_i) + \gamma \hat{V}_\theta^{\pi_\phi}(s_i) - \hat{V}_\theta^{\pi_\phi}(s_i)$
5. $\nabla_\phi J(\phi) \approx \sum_{i=1}^N \nabla_\phi \log \pi_\phi(a_i | s_i) \hat{A}^{\pi_\phi}(s_i, a_i)$
6. $\phi \leftarrow \phi + \alpha \nabla_\phi J(\phi)$



Off-policy AC

> Version 2. Replay buffer

1. collect experience $\{s_i, a_i\}$ from π_ϕ and add to replay buffer
 2. sample a batch $\{s_i, a_i, r_i, s_i'\}$ from buffer \mathcal{R}
 3. update \hat{Q}_θ^π using targets $r(s_i, a_i) + \gamma \hat{Q}_\theta^\pi(s_i', a_i')$ for each s_i, a_i
 4. evaluate $\hat{A}^{\pi_\phi}(s_i, a_i) = \hat{Q}_\theta^\pi(s_i, a_i) - \hat{V}_\theta^{\pi_\phi}(s_i)$
 5. $\nabla_\phi J(\phi) \approx \sum_{i=1}^N \nabla_\phi \log \pi_\phi(a_i | s_i) \hat{A}^{\pi_\phi}(s_i, a_i)$
 6. $\phi \leftarrow \phi + \alpha \nabla_\phi J(\phi)$
- $\hat{V}_\theta^\pi(s_i') = \hat{Q}_\theta^\pi(s_i', a_i'), a_i' \sim \pi_\phi$
not from replay buffer \mathcal{R}



Off-policy AC

> Version 2. Replay buffer

1. collect experience $\{s_i, a_i\}$ from π_ϕ and add to replay buffer
2. sample a batch $\{s_i, a_i, r_i, s_i'\}$ from buffer \mathcal{R}
3. update \hat{Q}_θ^π using targets $r(s_i, a_i) + \gamma \hat{Q}_\theta^\pi(s_i', a_i')$ for each s_i, a_i
4. evaluate $\hat{A}^{\pi_\phi}(s_i, a_i) = \hat{Q}_\theta^\pi(s_i, a_i) - \hat{V}_\theta^{\pi_\phi}(s_i)$
5. $\nabla_\phi J(\phi) \approx \sum_{i=1}^N \nabla_\phi \log \pi_\phi(a_i | s_i) \hat{A}^{\pi_\phi}(s_i, a_i)$
6. $\phi \leftarrow \phi + \alpha \nabla_\phi J(\phi)$

$$\nabla_\phi \log \pi_\phi(a_i | s_i) \hat{A}^{\pi_\phi}(s_i, a_i)$$

a_i is not from π_ϕ

sample $a_i^\pi \sim \pi_\phi(a | s_i)$
 $\nabla_\phi \log \pi_\phi(a_i^\pi | s_i) \hat{A}^{\pi_\phi}(s_i, a_i^\pi)$
in practice: use without baseline
 $\nabla_\phi \log \pi_\phi(a_i^\pi | s_i) \hat{Q}^{\pi_\phi}(s_i, a_i^\pi)$
(higher variance, but convenient)

Off-policy AC


- > Version 2. Replay buffer
 1. collect experience $\{s_i, a_i\}$ from π_ϕ and add to replay buffer
 2. sample a batch $\{s_i, a_i, r_i, s'_i\}$ from buffer \mathcal{R}
 3. update \hat{Q}_θ^π using targets $r(s_i, a_i) + \gamma \hat{Q}_\theta^\pi(s'_i, a'_i)$ for each s_i, a_i
 4. $\nabla_\phi J(\phi) \approx \sum_{i=1}^N \nabla_\phi \log \pi_\phi(a_i | s_i) \hat{Q}_\theta^{\pi_\phi}(s_i, a_i)$ where $a_i \sim \pi_\phi(a | s_i)$
 5. $\phi \leftarrow \phi + \alpha \nabla_\phi J(\phi)$
- > Any remaining problems?
 - s_i didn't come from $p_\phi(s)$
 - nothing we can do here, just accept it
 - intuition: we want optimal policy on $p_\phi(s)$, but we get optimal policy on a broader distribution

Off-policy AC


- > Version 2. Replay buffer
 1. collect experience $\{s_i, a_i\}$ from π_ϕ and add to replay buffer
 2. sample a batch $\{s_i, a_i, r_i, s'_i\}$ from buffer \mathcal{R}
 3. update \hat{Q}_θ^π using targets $r(s_i, a_i) + \gamma \hat{Q}_\theta^\pi(s'_i, a'_i)$ for each s_i, a_i
 4. $\nabla_\phi J(\phi) \approx \sum_{i=1}^N \nabla_\phi \log \pi_\phi(a_i | s_i) \hat{Q}_\theta^{\pi_\phi}(s_i, a_i)$ where $a_i \sim \pi_\phi(a | s_i)$
 5. $\phi \leftarrow \phi + \alpha \nabla_\phi J(\phi)$
- > Implementation details
 - How to fit Q-functions better?
 - To estimate the gradient better, use reparameterization trick
- > Example practical algorithm:
 - Soft Actor-critic: off-policy maximum entropy deep RL with a stochastic actor (Haarnoja, et al., 2018)

Off-policy AC

- > Off-policy with replay buffer (SAC) can be far more data efficient
- > They can also generally be a lot harder to tune hyperparameters, less stable (than PPO)




(e) Humanoid-v1




(f) Humanoid (rllab)

Off-policy AC


> Learning curves of algorithms (no free lunch theorem)




(b) HalfCheetah



(d) HumanoidStandup




(c) Hopper



(j) Walker2d

Off-policy AC

> Continuous state space



General RL tips

- > Start with the simple and fast algorithms
- > DQN, DDPG and PPO exhibit good performance overall
 - DQN, DDPG are easy to tune
 - PPO is harder to tune, but stable
 - can perform well with smaller model
 - good choice for cheap envs.
- > TD3 and SAC are good choices for expensive envs.
 - SAC generates stochastic policy that can be useful for exploration
 - harder to tune but can be more memory efficient
 - both benefit from larger network architectures
 - sensitive to input normalization
- > Even with the same algorithm, performance may vary significantly depending on the implementation details! (Use the verified libraries)

Trade-off between algorithms

- > Sampling efficiency
 - this is only a general trend
 - more efficient doesn't necessarily mean shorter wall clock time

