

ECE7121 Learning-based control – 2025 Fall

Optimal control




INHA UNIVERSITY

Overview

- > Feedback control
- > Optimization problem
- > Linear system and LQR
- > Trajectory optimization
- > Dynamic programming


Feedback control

- > Tracking a reference signal
 - Control input: steering wheel angle (θ)
 - Control output: vehicle yaw rate ($\dot{\psi}$)



Feedback control


- > Reference tracking with uncertainty



Feedback control


- > Feedback: compare measurement with a desired value and use the difference to determine the control action
- > Desiderata
 - stability: multiple notions; loosely system output is under control
 - tracking: the output should track the reference as closely as possible
 - disturbance rejection: the output should be insensitive to disturbance
 - robustness: controller should perform well up to some degree of model misspecification

The Map of Control Theory




The map of control theory

- > Intertwining optimal control and learning-based control to understand the framework and context



The map of control theory

- > Intertwining optimal control and learning-based control to understand the framework and context



Problem formulation

- > Mathematical description of the system to be controlled

- state variables: $x_1(t), x_2(t), \dots, x_n(t)$
- control inputs: $u_1(t), u_2(t), \dots, u_m(t)$
- model:

$$\begin{aligned}\dot{x}_1(t) &= f_1(x_1(t), x_2(t), \dots, x_n(t), u_1(t), u_2(t), \dots, u_m(t), t) \\ \dot{x}_2(t) &= f_2(x_1(t), x_2(t), \dots, x_n(t), u_1(t), u_2(t), \dots, u_m(t), t) \\ &\dots \\ \dot{x}_n(t) &= f_n(x_1(t), x_2(t), \dots, x_n(t), u_1(t), u_2(t), \dots, u_m(t), t)\end{aligned}$$

- In compact form: $\dot{\mathbf{x}}(t) = \mathbf{f}(\mathbf{x}(t), \mathbf{u}(t), t)$

Problem formulation

> Constraints

- initial and final conditions (boundary conditions)

$$x(t_0) = x_0, \quad x(t_f) = x_f$$

- constraints on state trajectories

$$\underline{X} \leq x(t) \leq \overline{X}$$

- control authority

$$\underline{U} \leq u(t) \leq \overline{U}$$

- and many more ...

Problem formulation

- > Performance measure

$$J = h(\mathbf{x}(t_f), t_f) + \int_{t_0}^{t_f} g(\mathbf{x}(t), \mathbf{u}(t), t) dt$$

h (terminal cost) and g (running cost) are scalar functions

- > Problem: find an admissible control \mathbf{u}^* which causes the system

$$\dot{\mathbf{x}}(t) = \mathbf{f}(\mathbf{x}(t), \mathbf{u}(t), t)$$

to follow an admissible trajectory \mathbf{x}^* that minimizes the performance measure

Optimization problem

- > Finding the minimizer of a function subject to constraints:

$$\underset{x}{\text{minimize}} J(x)$$

$$\begin{aligned} \text{s.t. } & f_i(x) \leq 0, \quad i = \{1, \dots, k\} \\ & h_j(x) = 0, \quad j = \{1, \dots, l\} \end{aligned}$$

> Gradient: $\nabla f = \begin{bmatrix} \frac{\partial f}{\partial x_1} \\ \vdots \\ \frac{\partial f}{\partial x_n} \end{bmatrix}$, Hessian: $\nabla^2 f(x) = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \frac{\partial^2 f}{\partial x_n \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_n^2} \end{bmatrix}$

Optimization problem

> Unconstrained optimization


- minimize $f(x)$
- Necessary condition: $\nabla f(x) = 0, \nabla^2 f(x^*) \geq 0$
- Sufficient condition: $\nabla f(x) = 0, \nabla^2 f(x^*) > 0$

> Iterative gradient descent

- $x_{t+1} = x_t - \alpha_t \nabla f(x_t)$

> Newton's method


- $x_{t+1} = x_t - (\nabla^2 f(x_t))^{-1} \nabla f(x_t)$



Optimization problem

> Optimization with equality constraints

- minimize $f(x)$
s. t. $\stackrel{x}{h_i(x)} = 0, \quad i = 1, \dots, m$



Optimization problem

> Lagrange multipliers

- for a given local minimum \mathbf{x}^* , there exists scalars $\lambda_1, \dots, \lambda_m$ called Lagrange multipliers such that

$$\nabla f(\mathbf{x}^*) + \sum_{i=1}^m \lambda_i \nabla h_i(\mathbf{x}^*) = 0$$

- If we had no constraints, the solution satisfies $\nabla f(\mathbf{x}^*) = 0$
- With equality constraints, we are not allowed to move in all directions
- We can move only along directions that stay on the constraint surface
- At an optimum, any movement along the constraint surface should not decrease the value of the objective. for direction d , $\nabla h_i(\mathbf{x}^*)^\top d = 0$
- If we move in a direction d , directional derivative of f along d must be zero (otherwise you could decrease f), $\nabla f(\mathbf{x}^*)^\top d = 0$
- The gradient of f must be orthogonal to the space of feasible directions
- $\nabla f(\mathbf{x}^*)$ must lie in the span of $\nabla h_i(\mathbf{x}^*)$, $\nabla f(\mathbf{x}^*) = \sum_{i=1}^m \lambda_i \nabla h_i(\mathbf{x}^*)$


Optimization problem

> Lagrange multipliers


- for a given local minimum \mathbf{x}^* , there exists scalars $\lambda_1, \dots, \lambda_m$ called Lagrange multipliers such that

$$\nabla f(\mathbf{x}^*) + \sum_{i=1}^m \lambda_i \nabla h_i(\mathbf{x}^*) = 0$$

$$f(\mathbf{x}) = x_1 + x_2$$



$$h(\mathbf{x}) = x_1^2 + x_2^2 - 2$$




Optimization problem

> Lagrange multipliers


- for a given local minimum \mathbf{x}^* , there exists scalars $\lambda_1, \dots, \lambda_m$ called Lagrange multipliers such that

$$\nabla f(\mathbf{x}^*) + \sum_{i=1}^m \lambda_i \nabla h_i(\mathbf{x}^*) = 0$$

$$f(\mathbf{x}) = x_1 + x_2$$



$$h(\mathbf{x}) = x_1^2 + x_2^2 - 2$$



Optimization problem

> Lagrange multipliers

- for a given local minimum \mathbf{x}^* , there exists scalars $\lambda_1, \dots, \lambda_m$ called Lagrange multipliers such that

$$\nabla f(\mathbf{x}^*) + \sum_{i=1}^m \lambda_i \nabla h_i(\mathbf{x}^*) = 0$$

$$f(\mathbf{x}) = x_1 + x_2 \quad h(\mathbf{x}) = x_1^2 + x_2^2 - 2$$

- $\begin{bmatrix} 1 \\ 1 \end{bmatrix} = \lambda \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}, \quad x_1^2 + x_2^2 - 2 = 0$


$$\frac{2}{\lambda^2} = 2, \rightarrow \lambda = \pm 1 \rightarrow (x_1, x_2) = (\pm 1, \pm 1)$$

- Optimal point: $(x_1^*, x_2^*) = (-1, -1)$

Optimization problem

- > Optimization with inequality constraints

- minimize $f(x)$
s. t. $\begin{matrix} x \\ g_i(x) \geq 0, \quad i = 1, \dots, r \end{matrix}$




Optimization problem

- > It is similar to equality constrained optimization

- $g(x^*) = 0$
- $\nabla f(x^*) + \sum_{i=1}^r \mu_i \nabla g_i(x^*) = 0$

Case 1: x^* is on the boundary of the feasible region.




if the constraint $g_i(x) \geq 0, \mu_i \leq 0$
if the constraint $g_i(x) \leq 0, \mu_i \geq 0$

Optimization problem

- > It is similar to unconstrained optimization
 - $g(x^*) > 0$
 - $\nabla f(x^*) = 0$
 - equivalent to $\nabla f(x^*) + \sum_{i=1}^r \mu_i \nabla g_i(x^*) = 0, \mu_i = 0$

Case 2: x^* is inside the feasible region.



Optimization problem

> minimize $f(\mathbf{x})$

$$\begin{aligned} & \underset{\mathbf{x}}{\text{s.t.}} h_i(\mathbf{x}) = 0, \quad i = 1, \dots, m \\ & \qquad g_i(\mathbf{x}) \leq 0, \quad i = 1, \dots, r \end{aligned}$$

> Define Lagrange function

$$L(\mathbf{x}, \boldsymbol{\lambda}, \boldsymbol{\mu}) = f(\mathbf{x}) + \sum_{i=1}^m \lambda_i h_i(\mathbf{x}) + \sum_{j=1}^r \mu_j g_j(\mathbf{x})$$

$$\left. \begin{array}{l} \text{- } \nabla_{\mathbf{x}} L(\mathbf{x}, \boldsymbol{\lambda}, \boldsymbol{\mu}) = 0 \\ \text{- } \mu_i \geq 0 \\ \text{- } \mu_i^* g_i(\mathbf{x}^*) = 0 \\ \text{- } h_i(\mathbf{x}) = 0 \\ \text{- } g_i(\mathbf{x}) \leq 0 \end{array} \right\} \rightarrow \nabla f(\mathbf{x}) + \sum_{i=1}^m \lambda_i \nabla h_i(\mathbf{x}) + \sum_{j=1}^r \mu_j \nabla g_j(\mathbf{x}) = 0$$

Necessary optimality condition called
Karush-Kuhn-Tucker (KKT) condition

Linear dynamical system

> Model: $\mathbf{x}_{t+1} = A_t \mathbf{x}_t + B_t \mathbf{u}_t, \quad \mathbf{y}_t = C_t \mathbf{x}_t$

> Linear quadratic control

- minimize $\underset{\mathbf{u}}{J_y} + \rho J_{\mathbf{u}}$

$$\text{s.t. } \mathbf{x}_{t+1} = A_t \mathbf{x}_t + B_t \mathbf{u}_t, \quad t = 1, \dots, T-1 \\ \mathbf{x}_1 = \mathbf{x}_{init}, \quad \mathbf{x}_T = \mathbf{x}_{des}$$

$$J_y = \sum_{t=1}^T \|\mathbf{y}_t\|^2 = \|C\mathbf{x}_t\|^2, \quad J_{\mathbf{u}} = \sum_{t=1}^T \|\mathbf{u}_t\|^2$$

- first constraint imposes the linear dynamics equations
- second set of constraints specifies the initial and final state
- \mathbf{y} would be the error state (the difference from the desired state)

Linear dynamical system

- > Can be written as

$$\begin{array}{ll}\text{minimize}_{\mathbf{z}} & \|\tilde{\mathbf{A}}\mathbf{z} - \tilde{\mathbf{b}}\|^2 \\ \text{s.t.} & \tilde{\mathbf{C}}\mathbf{z} = \tilde{\mathbf{d}}\end{array}$$

- > KKT conditions in matrix form

- $\begin{bmatrix} 2\tilde{\mathbf{A}}^\top \tilde{\mathbf{A}} & \tilde{\mathbf{C}}^\top \\ \tilde{\mathbf{C}} & 0 \end{bmatrix} \begin{bmatrix} \mathbf{z} \\ \lambda \end{bmatrix} = \begin{bmatrix} 2\tilde{\mathbf{A}}^\top \tilde{\mathbf{b}} \\ \tilde{\mathbf{d}} \end{bmatrix}$

- > Solve the KKT conditions to find candidate solutions
then, evaluate the original cost function to determine the optimum

$$\tilde{\mathbf{A}} = \left[\begin{array}{ccc|ccc} C_1 & \cdots & 0 & 0 & \cdots & 0 \\ \vdots & \ddots & \vdots & \vdots & \ddots & \vdots \\ 0 & \cdots & C_T & 0 & \cdots & 0 \\ \hline 0 & \cdots & 0 & \sqrt{\rho}I & \cdots & 0 \\ \vdots & \ddots & \vdots & \vdots & \ddots & \vdots \\ 0 & \cdots & 0 & 0 & \cdots & \sqrt{\rho}I \end{array} \right], \quad \tilde{\mathbf{b}} = 0$$

$$\tilde{\mathbf{C}} = \left[\begin{array}{ccccc|ccccc} A_1 & -I & 0 & \cdots & 0 & 0 & B_1 & 0 & \cdots & 0 \\ 0 & A_2 & -I & \cdots & 0 & 0 & 0 & B_2 & \cdots & 0 \\ \vdots & \vdots & \vdots & & \vdots & \vdots & \vdots & \vdots & & \vdots \\ 0 & 0 & 0 & \cdots & A_{T-1} & -I & 0 & 0 & \cdots & B_{T-1} \\ \hline I & 0 & 0 & \cdots & 0 & 0 & 0 & 0 & \cdots & 0 \\ 0 & 0 & 0 & \cdots & 0 & I & 0 & 0 & \cdots & 0 \end{array} \right], \quad \tilde{\mathbf{d}} = \begin{bmatrix} 0 \\ 0 \\ \vdots \\ 0 \\ x^{\text{init}} \\ x^{\text{des}} \end{bmatrix}$$

Linear dynamical system

- > Linear quadratic regulator (LQR)
 - with linear dynamics model: $\mathbf{x}_{t+1} = A_t \mathbf{x}_t + B_t \mathbf{u}_t$ and quadratic cost f.
 - minimize $\frac{1}{2} \mathbf{x}_T^\top L \mathbf{x}_T + \sum_{t=0}^{T-1} \frac{1}{2} \mathbf{x}_t^\top Q \mathbf{x}_t + \frac{1}{2} \mathbf{u}_k^\top R \mathbf{u}_k$
s.t. $\mathbf{x}_{t+1} = A_t \mathbf{x}_t + B_t \mathbf{u}_t, \quad t = 1, \dots, T - 1$
 $\mathbf{x}_1 = \mathbf{x}_{init}, \quad \mathbf{x}_T = \mathbf{x}_{des}$
 - L, Q, R are positive definite weight matrices
 - We want to stabilize the system in the end, so we prioritize the terminal state

Linear dynamical system

> Constrained LQR

- minimize $\frac{1}{2} \mathbf{x}_T^\top L \mathbf{x}_T + \sum_{t=0}^{T-1} \frac{1}{2} \mathbf{x}_t^\top Q \mathbf{x}_t + \frac{1}{2} \mathbf{u}_k^\top R \mathbf{u}_k$
s.t. $\mathbf{x}_{t+1} = A_t \mathbf{x}_t + B_t \mathbf{u}_t, \quad t = 1, \dots, T - 1$
 $\mathbf{x}_1 = \mathbf{x}_{init}, \quad \mathbf{x}_T = \mathbf{x}_{des}$
 $\mathbf{u}_{\min} \leq \mathbf{u}_t \leq \mathbf{u}_{\max}, \quad \mathbf{x}_{\min} \leq \mathbf{x} \leq \mathbf{x}_{\max}$ (inequality constraints)

- In reality, there are always control input and state boundaries

> It is a special case of quadratic program (QP)


- minimize $\frac{1}{2} \mathbf{x}^\top Q \mathbf{x} + q \mathbf{x}$
s.t. $A_{in} \mathbf{x} \leq b_{in}$ Q is positive definite.
 $A_{eq} \mathbf{x} = b_{eq}$
 $lb \leq \mathbf{x} \leq ub$
- \mathbf{x}^* exists and is unique (QP is convex and feasible) unless non-empty constraint subset

Convex optimization

- > With convex cost function, convex inequality constraints, and affine equality constraints


Nonconvex Optimization

- > Multiple local minima ⭐
- > Sensitive to initial point ●



Convex Optimization

- > Unique minimum: global/local




Linear dynamical system

- > LQR solution gives all the sequences of state and input
- > Applying the control input sequence leads to an open-loop control policy
 - there is no feedback along the way
 - LQR is agnostic to future changes during execution
 - it can't handle the model error and estimation error
- > Instead of executing everything, just execute the first action and do the optimization again = MPC
 - it needs quick optimization (we need it every timestep)
 - we have a very complicated feedback, it is difficult to analyze it
 - still, works very well in practice

Model predictive control

- > With finite horizon, iterative planning framework



- Depending on the optimization solver
 - nonlinear program (nonlinear solver) = nonlinear MPC
 - quadratic program = QP-MPC = linear MPC (recommended!)
 - sampling-based optimization >> model-based RL


Trajectory optimization

- > TO is a collection of techniques that are used to find open-loop solutions to an optimal control problem


- minimize $\underset{\boldsymbol{u}}{N(\boldsymbol{x}_T)} + \int_0^T L(\boldsymbol{x}_t, \boldsymbol{u}_t) dt$
s.t. $\dot{\boldsymbol{x}} = f(\boldsymbol{x}_t, \boldsymbol{u}_t, t)$
 $\boldsymbol{x}(0) = \boldsymbol{x}_{init}, \boldsymbol{x}(T) = \boldsymbol{x}_{des}$
+ additional constraints

- additional constraints include path constraints and bounds on $\boldsymbol{x}, \boldsymbol{u}$

Open-Loop Solution (optimal trajectory)



Closed-Loop Solution (optimal policy)



Trajectory optimization

- > Indirect methods (optimize -> discretize)
 - analytically construct the conditions for optimality
 - Euler-Lagrange eqn., Hamilton-Jacobi-Bellman eqn.
 - more accurate, harder to pose and solve
 - usually used in aerospace industry
- > Direct methods (discretize -> optimize)
 - numerically optimizing it as a nonlinear program
 - more recent and more widely used in robotics
 - less accurate, easier to pose and solve
- > Solving strategies
 - shooting: only solve for inputs and simulate the state
 - collocation: parametrize the problem via many trajectory segments\
 - These are connected to model-based RL in future

Dynamic programming


> Focus is now on finding optimal closed-loop policies

- cost: $J = h(\mathbf{x}(t_f), t_f) + \int_{t_0}^{t_f} g(\mathbf{x}(t), \mathbf{u}(t), t) dt$
- model: $\dot{\mathbf{x}}(t) = \mathbf{f}(\mathbf{x}(t), \mathbf{u}(t), t)$
- control constraints: $\mathbf{u}(t) \in U(\mathbf{x}(t))$
- Find $\mathbf{u}^*(t) = \pi^*(\mathbf{x}_t)$

> Principle of optimality

- the key concept behind the DP is the principle of optimality
- cost a-e = cost a-b + cost b-e

$$J_{ae} = J_{ab} + J_{be}$$



Dynamic programming


- > If $a - b - e$ is optimal path from a to e , then $b - e$ is optimal path from b to e

- proof:

- suppose $b - c - e$ is the optimal path from b to e

- then, $J_{bce} < J_{be}$



- and $J_{ab} + J_{bce} < J_{ab} + J_{be} = J_{ae}^*$ (contradiction !)



- > Principle of optimality: tail of optimal sequences is optimal for tail subproblems

Dynamic programming

- > Applying the principle of optimality
 - the optimal trajectory is found by comparing
$$C_{bcf} = J_{bc} + J_{cf}^*$$
$$C_{bdf} = J_{bd} + J_{df}^*$$
$$C_{bef} = J_{be} + J_{ef}^*$$
 - carry out this procedure backward in time



Dynamic programming

- > In most cases, DP algorithm needs to be performed numerically
- > A few case can be solved analytically (LQR)

- minimize $\frac{1}{2} \mathbf{x}_T^\top L \mathbf{x}_T + \sum_{t=0}^{T-1} \frac{1}{2} \mathbf{x}_t^\top Q \mathbf{x}_t + \frac{1}{2} \mathbf{u}_k^\top R \mathbf{u}_k$

$$\text{s.t. } \mathbf{x}_{t+1} = A_t \mathbf{x}_t + B_t \mathbf{u}_t, \quad t = 1, \dots, T-1$$
$$\mathbf{x}_1 = \mathbf{x}_{init}, \quad \mathbf{x}_T = \mathbf{x}_{des}$$

- $J_T^*(\mathbf{x}_T) = \frac{1}{2} \mathbf{x}_T^\top L \mathbf{x}_T$

- going backward

$$J_{T-1}(\mathbf{x}_{T-1}) = \underset{\mathbf{u}_{T-1}}{\text{minimize}} \frac{1}{2} \mathbf{x}_{T-1}^\top Q \mathbf{x}_{T-1} + \frac{1}{2} \mathbf{u}_{T-1}^\top R \mathbf{u}_{T-1} + \frac{1}{2} \mathbf{x}_T^\top L \mathbf{x}_T$$
$$= \underset{\mathbf{u}_{T-1}}{\text{minimize}} \frac{1}{2} \mathbf{x}_{T-1}^\top Q \mathbf{x}_{T-1} + \frac{1}{2} \mathbf{u}_{T-1}^\top R \mathbf{u}_{T-1} +$$
$$\frac{1}{2} (A_{T-1} \mathbf{x}_{T-1} + B_{T-1} \mathbf{u}_{T-1})^\top H (A_{T-1} \mathbf{x}_{T-1} + B_{T-1} \mathbf{u}_{T-1})$$

Dynamic programming

- Taking derivative

$$\frac{\partial J_{T-1}^*(\mathbf{x}_{T-1})}{\partial \mathbf{u}_{T-1}} = R\mathbf{u}_{T-1} + B_{T-1}^\top H(A_{T-1}\mathbf{x}_{T-1} + B_{T-1}\mathbf{u}_{T-1}) = 0$$

$$\mathbf{u}_{T-1}^* = -(R + B_{T-1}^\top H B_{T-1})^{-1} B_{T-1}^\top H A_{T-1} \mathbf{x}_{T-1} = F_{T-1} \mathbf{x}_{T-1}$$

$$\begin{aligned} J_{T-1}(\mathbf{x}_{T-1}) &= \frac{1}{2} \mathbf{x}_{T-1}^\top \{Q + F_{T-1}^\top R F_{T-1} + \\ &\quad (A_{T-1} + B_{T-1} F_{T-1})^\top H (A_{T-1} + B_{T-1} F_{T-1})\} \mathbf{x}_{T-1} \\ &= \mathbf{x}_{T-1}^\top P_{T-1} \mathbf{x}_{T-1} \end{aligned}$$

- Proceeding by induction

- $\mathbf{u}_t^* = F_t \mathbf{x}_t$ where $F_t = -(R + B_t^\top P_{t+1} B_t)^{-1} B_t^\top P_{t+1} A_t$
- $J_t(\mathbf{x}_t) = \frac{1}{2} \mathbf{x}_t^\top P_t \mathbf{x}_t$ where $P_t = Q + F_t^\top R F_t + (A_t + B_t F_t)^\top P_{t+1} (A_t + B_t F_t)$
- At the end we could get $J_0(\mathbf{x}_0) = \frac{1}{2} \mathbf{x}_0^\top P_0 \mathbf{x}_0$

Dynamic programming

- > We can do same things in continuous-time setting
 - DP leads to Hamilton-Jacobi-Bellman (HJB), Hamilton-Jacobi-Issacs (HJI) equations
 - HJB: $\frac{\partial J}{\partial t} + \min_u [g(x, u, t) + \frac{\partial J}{\partial x} f(x, u, t)] = 0$
 - but, solving PDEs is hard except in simple cases
 - for high-dimensional systems, numerical solutions become infeasible

What if the problem is nonconvex?

- > Local optimization method
 - sequential quadratic/convex programming (SQP, SCP)
- > Dynamic programming + local optimization
 - Iterative linear quadratic regulator (iLQR)
 - Differential dynamic programming (DDP)
- > Sampling-based methods
 - Model predictive path integral control (MPPI)
 - Cross-entropy method (CEM)
 - Covariance matrix adaptation evolutionary strategy (CMA-ES)