

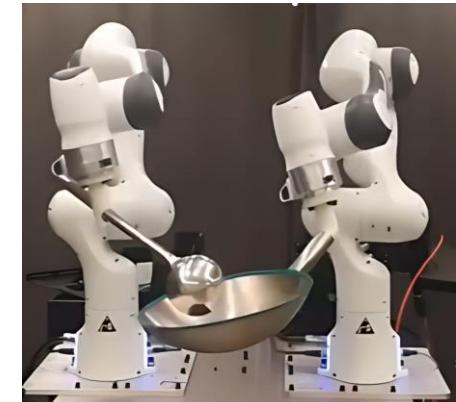
ECE7121 Learning-based control – 2025 Fall

Model-based RL



INHA UNIVERSITY

What's wrong with known dynamics?



Can we learn a simulator?

- > Simulator : predictive model of s_{t+1} given s_t, a_t
= transition dynamics $f(s_t, a_t) = s_{t+1}$
- > How to apply RL algorithm in the real-world?
 - collect data $D = \{(s, a, s')_i\}$ using some base policy
 - $f(s, a)$ to minimize $\sum_i \|f(s_i, a_i) - s'_i\|^2$
 - run your favorite RL method inside simulator f
- > Issues
 - Data coverage matters a lot
 - How to learn a simulator (model) efficiently?
 - Need to account for model inaccuracies



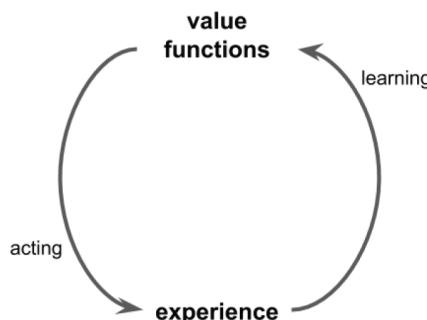
Can we learn a simulator?

> Simulators in brief

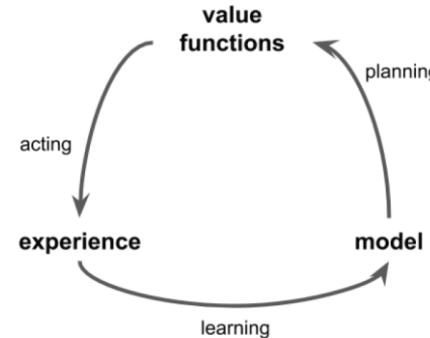
- you have it already
 - games (e.g., go, chess)
 - derived from the first principle (e.g., rigid body dynamics)
 - simulated environments (robot simulators)
- you approximately know most of it (e.g. certain physical models)
 - can fit the unknown parameters using the data – system identification
- you don't know it (most scenarios in practice)
 - learn it end-to-end
 - learn a (low-dim) state representation, then learn model over representations

Can we exploit a simulator?

- > If we knew $f(s_t, a_t) = s_{t+1}$, we could use the tools from last week
- > Model-based RL v.0.5
 - run base policy $\pi_0(a_t|s_t)$ (e.g., random policy) to collect $\mathcal{D} = \{(s, a, s')_i\}$
 - learn dynamics model $f(s, a)$ to minimize $\sum_i \|f(s_i, a_i) - s'_i\|^2$
 - plan through $f(s, a)$ to choose actions



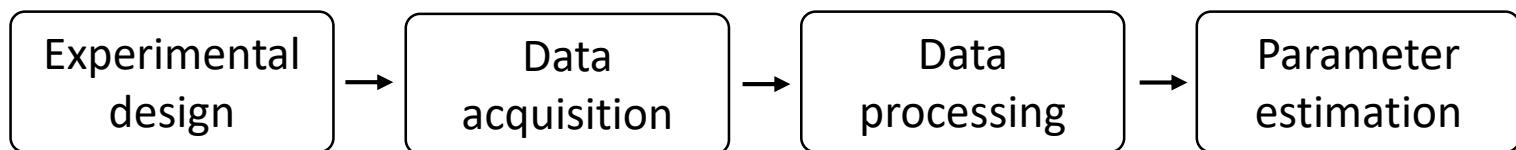
model-free RL



model-based RL

Can we exploit a simulator?

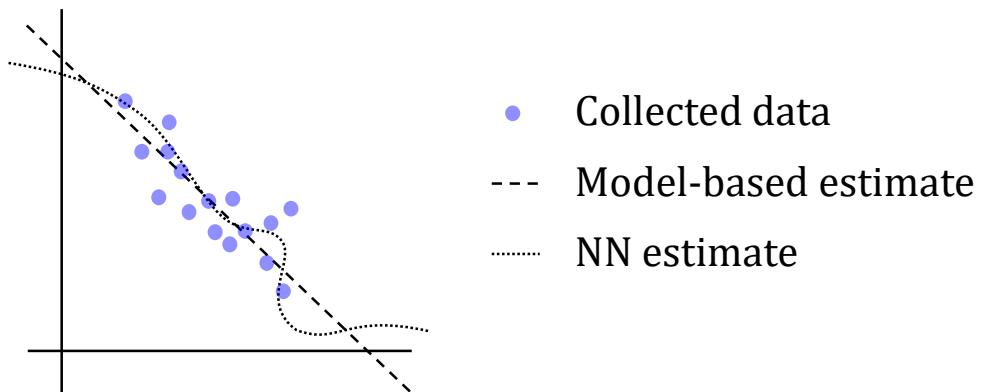
- > Model-based RL v.0.5, does it work well?
 - essentially how system identification works in classical robotics
 - some care should be taken to design a good base policy
 - particularly effective if we can hand-engineer a dynamics model using the domain knowledge, and fit just a few parameters
- > System identification procedure



Can we exploit a simulator?

- > Model-based RL v.0.5, does it work well?
 - not for the deep RL approaches
 - distribution mismatch problem becomes exacerbated as we use more expressive model classes

Linear regression $y = ax + b$

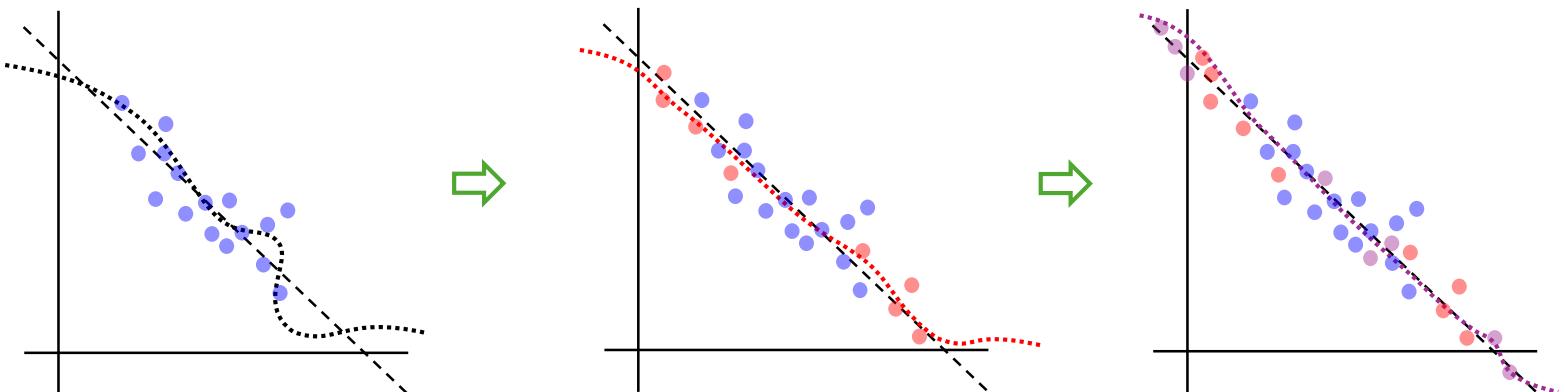


Model-based RL

> Model-based RL v.1.0

- run base policy $\pi_0(a_t|s_t)$ (e.g., random policy) to collect $\mathcal{D} = \{(s, a, s')_i\}$
- - learn dynamics model $f(s, a)$ to minimize $\sum_i \|f(s_i, a_i) - s'_i\|^2$
- plan through $f(s, a)$ to choose actions
- execute those actions and add the resulting data $\{(s, a, s')_i\}$ to \mathcal{D}

Linear regression $y = ax + b$



- Collected data
- Model-based estimate
- NN estimate

- Newly collected data

Model-based RL

> Model-based RL v.1.1

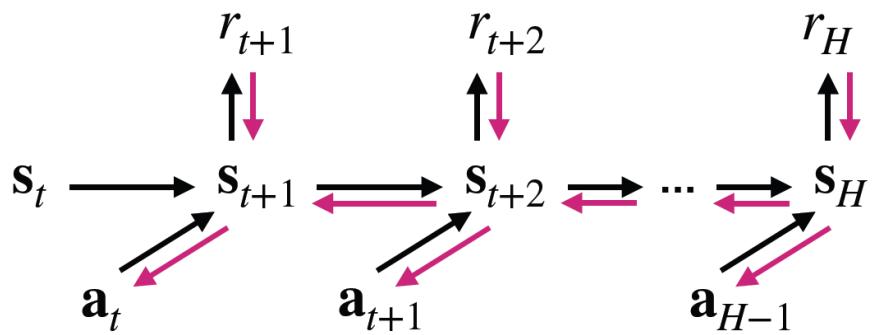
- run base policy $\pi_0(a_t|s_t)$ (e.g., random policy) to collect $\mathcal{D} = \{(s, a, s')_i\}$
- - learn dynamics model $f(s, a)$ to minimize $\sum_i \|f(s_i, a_i) - s'_i\|^2$
- - plan through $f(s, a)$ to choose actions
- - execute the **first action** and add the resulting data $\{(s, a, s')_i\}$ to \mathcal{D}

> We use MPC framework for planning

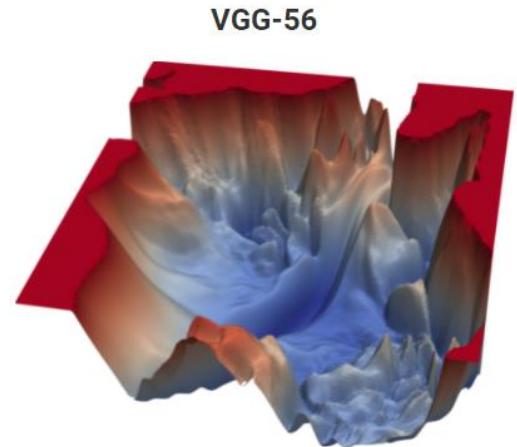
- Replanning helps with model errors
- The more you replan, the less perfect each individual plan needs to be
- but computation become intensive (appropriate for short horizons)

Planning with the learned model

- > Plan through $f(s, a)$ to choose actions (= trajectory optimization)
- > 1) gradient-based optimization



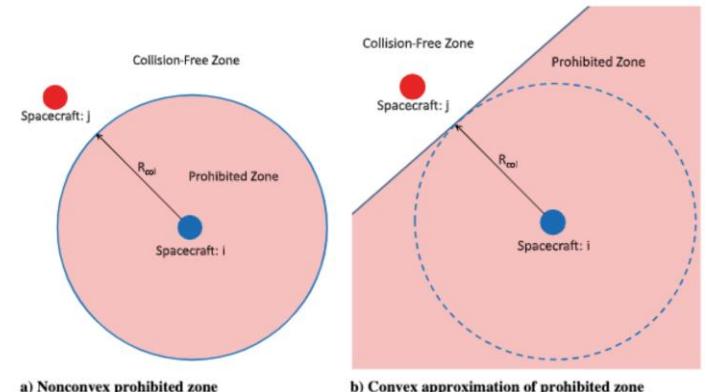
- It is a highly non-convex problem
- backpropagation uses the chain rule to compute gradients
- iterative gradient descent (basic, not efficient)
- only possible to differential simulators



Planning with the learned model

- > 1) gradient-based optimization
 - iterative convexification and optimization (SCP/SQP)
 - SQP Algorithm
 - Initialize a trajectory
 - affinized the dynamics
 - quadratic approximation of the cost
 - affinized the constraint
 - solve the QP

until convergence



- SCP uses arbitrary convex approximations (more general)
- SCP and SQP are heuristic
 - it can fail to find an optimal (feasible) solution
 - highly dependent on initialization

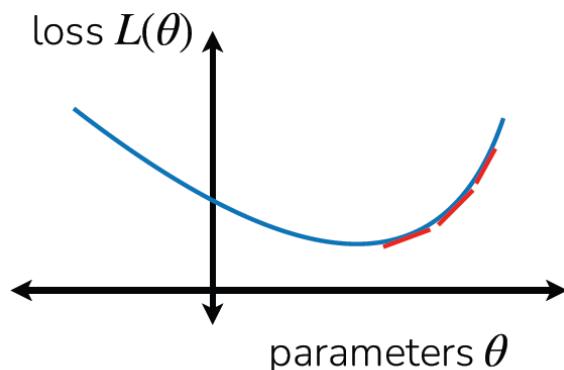
Planning with the learned model

- > 1) gradient-based optimization
 - iterative linear quadratic regulator (iLQR)
 - linearize the dynamics and makes a quadratic approximation of cost
 - solve using DP (Riccati equation)
 - don't directly deal with constraints (use penalty methods)
 - Differential dynamic programming (DDP)
 - use a full second-order method (need a second-order approximation of dyn.)
 - Both can be viewed as (approximated) Newton's method
 - Compare to QP, complexity increases linearly as horizon increases

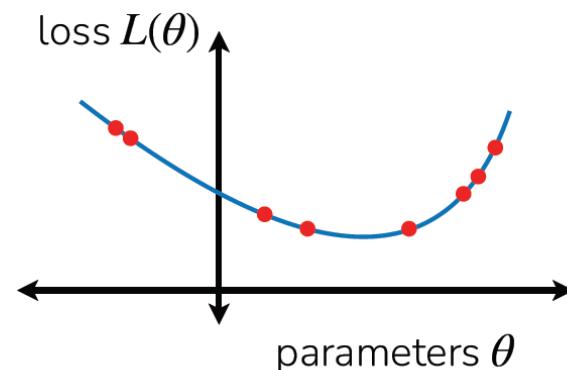
Planning with the learned model

> 2) sampling-based optimization

Gradient-based (1st order)



Sampling-based (0th order)



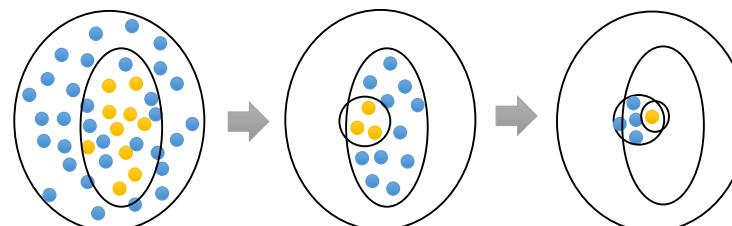
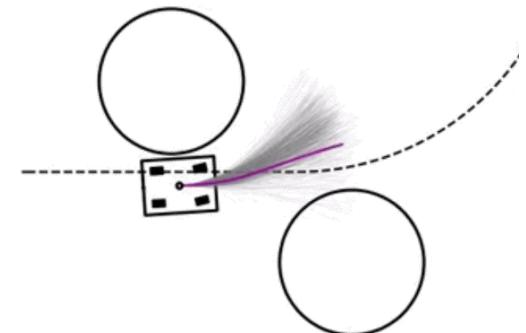
- requires no gradient information (gradient-free optimization)
- use randomness to explore the space
- often better at escaping local minima than gradient-based methods
- parallelizable
- scales poorly to high dimensions

Planning with the learned model

> 2) sampling-based optimization

- Random shooting
 - Guess and check
 - sample many action sequences ($A = \{a_t, \dots, a_{t+H}\}$) A_1, \dots, A_n from some distribution (e.g., uniform)
 - choose the best A_i (max return)
- can we improve the sampling distribution? (prior knowledge)
- Cross-entropy method
 - - sample many action sequences A_1, \dots, A_n from $p(A)$
 - evaluate and pick elites A_{i_1}, \dots, A_{i_m}
 - refit $p(A)$ to the elites A_{i_1}, \dots, A_{i_m}

until convergence



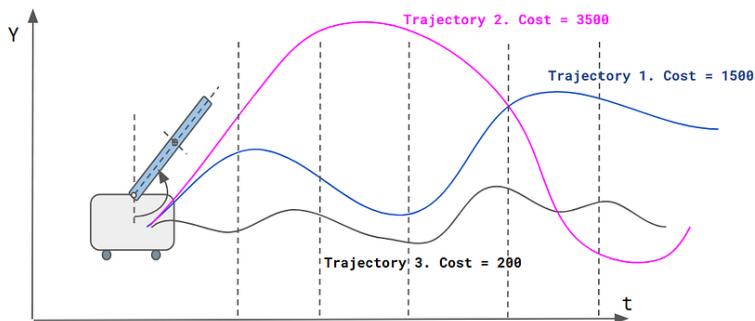
Planning with the learned model

> 2) sampling-based optimization

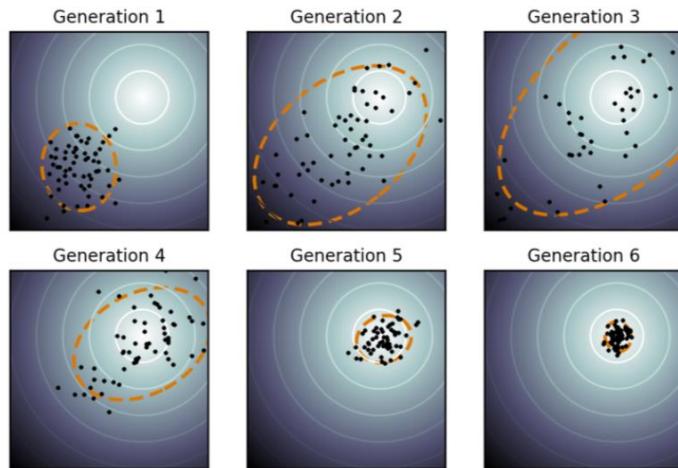
- model predictive path integral control (MPPI)
 - - sample many action sequences A_1, \dots, A_n from $p(A) \sim \mathcal{N}(\mu, \Sigma)$
 - evaluate corresponding costs C_1, \dots, C_n
 - compute weights based on costs: $w_i = \frac{\exp(-\frac{1}{\lambda}C_i)}{\sum_j \exp(-\frac{1}{\lambda}C_j)}$
 - update $\mu \leftarrow \sum_i w_i A_i$
- covariance matrix adaptation evolutionary strategy (CMA-ES)
 - - sample many action sequences A_1, \dots, A_n from $p(A) \sim \mathcal{N}(\mu, \Sigma)$
 - evaluate corresponding costs C_1, \dots, C_n
 - compute weights based on costs: w_i (e.g., exponential, top-K, ...)
 - update $\mu \leftarrow (1 - \gamma_\mu)\mu + \gamma_\mu \sum_i w_i A_i, \quad \Sigma \leftarrow (1 - \gamma_\Sigma)\Sigma + \gamma_\Sigma \sum_i (A_i - \mu)(A_i - \mu)^\top$
- Even more general: not necessarily Gaussian, parametrized distribution

Planning with the learned model

- > 2) sampling-based optimization
 - model predictive path integral control (MPPI)

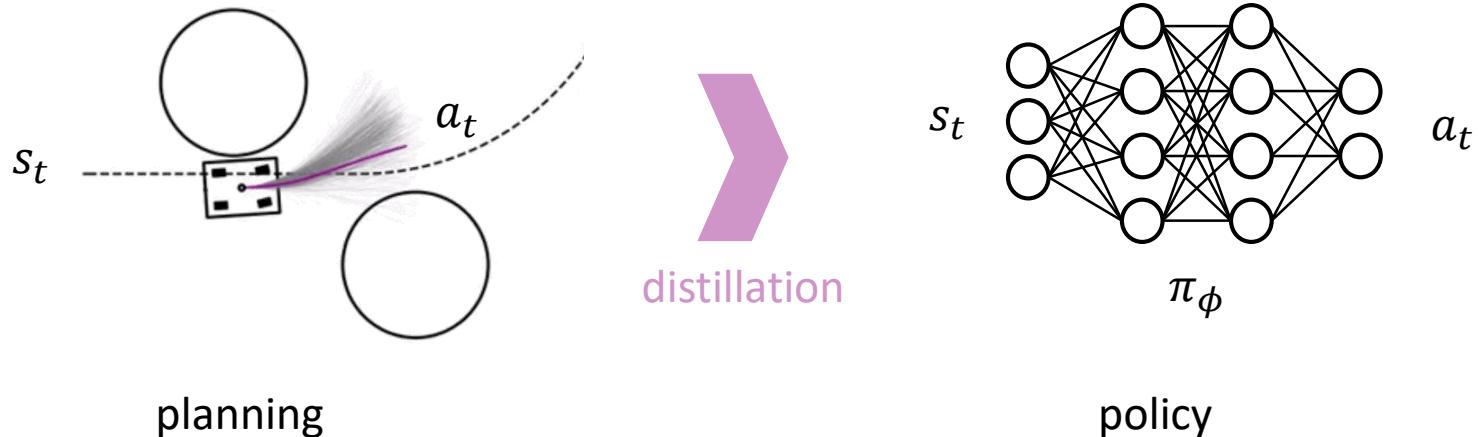


- covariance matrix adaptation evolutionary strategy (CMA-ES)



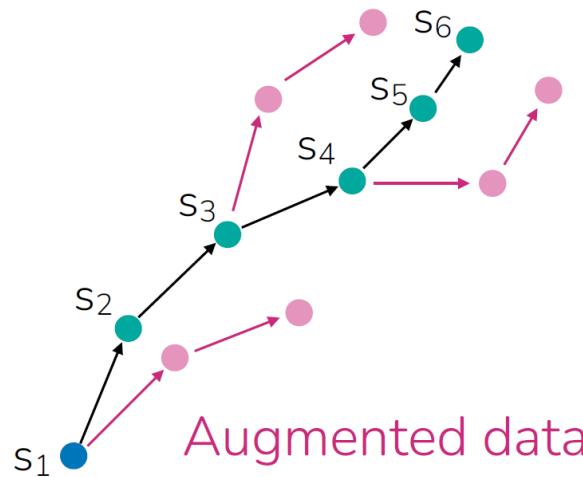
Planning with the learned model

- > Planning at test time can be computationally intensive
- > Option 1: distill planner's actions into a policy
 - still limited to short-horizon problems



Model-based policy optimization

- > Option 2: augment model-free RL methods with data from model
 - generate **full trajectories** from initial states?
 - model may not be accurate for long horizons
 - generate partial trajectories from initial states?
 - model may not get good coverage of later states
 - generate **partial trajectories** from all states in the data
 - add augmented dataset and update policy



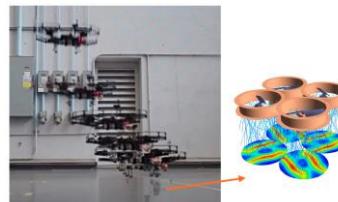
Model learning

- > Model-based RL v.0.5
 - run **base policy** $\pi_0(a_t|s_t)$ (e.g., random policy) to collect $\mathcal{D} = \{(s, a, s')_i\}$
 - learn **dynamics model** $f(s, a)$ to minimize $\sum_i \|f(s_i, a_i) - s'_i\|^2$
 - plan through $f(s, a)$ to choose actions
- > Issues
 - (1) Domain shift, over-fitting, and generalization
 - what policy we should use to generate data?
 - (2) good model might be bad for control
 - what function class we should use to learn?
 - (3) real-world models are not time-invariant
 - (4) high-dimensional state/action or rich observation settings
 - (5) under-fitting and compounding errors

Regularized model learning (1), (2)

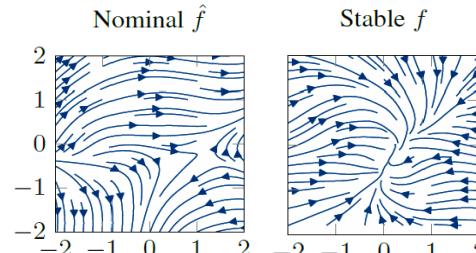
- > Not only minimize the regression error, also design regularization
 - f_θ has some good properties for control (Lipschitz, smooth, controllable, ...)
 - f_θ matches the physics (Lagrangian systems, symmetric, invariance ...)
 - $\sum_i \|f_\theta(s_i, a_i) - s'_i\|^2 + R(\theta)$ or $\sum_i \|f_\theta(s_i, a_i) - s'_i\|^2, s.t \theta \in \Theta$
 - Examples
 - Residual dynamics learning use spectral normalization $\|f_\theta(x, u)\|_{Lip} \leq \gamma$

Neural Lander: Stable Drone Landing Control Using Learned Dynamics
Guanya Shi^{*1}, Xichen Shi^{*1}, Michael O'Connell^{*1}, Rose Yu², Kamyar Azizzadenesheli³,
Animashree Anandkumar¹, Yisong Yue¹, and Soon-Jo Chung¹



- Learning stabilizable nonlinear dynamics

Learning Stabilizable Nonlinear Dynamics
with Contraction-Based Regularization
Sumeet Singh¹, Spencer M. Richards¹, Vikas Sindhwani², Jean-Jacques E. Slotine³, and
Marco Pavone¹



Regularized model learning (1), (2)

> Examples

- deep Lagrangian networks

LAGRANGIAN NEURAL NETWORKS

Miles Cranmer
Princeton University
mcranmer
@princeton.edu

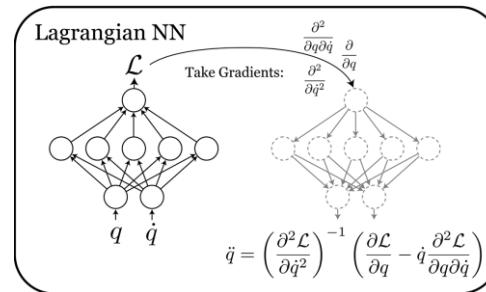
Sam Greydanus
Oregon State University
greydanus.17
@gmail.com

Stephan Hoyer
Google Research
shoyer
@google.com

Peter Battaglia
DeepMind
peterbattaglia
@google.com

David Spergel^{*}
Flatiron Institute
davidspergel
@flatironinstitute.org

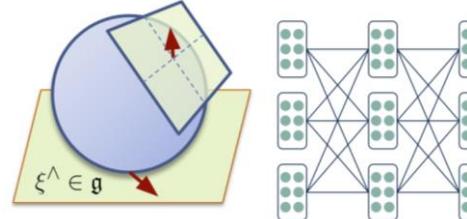
Shirley Ho[†]
Flatiron Institute
shirleyho
@flatironinstitute.org



- Invariance and equivariance (symmetry)

Lie Neurons: Adjoint-Equivariant Neural Networks for Semisimple Lie Algebras

Tzu-Yuan Lin ^{*1} Minghan Zhu ^{*1} Maani Ghaffari ¹



- Deep Koopman

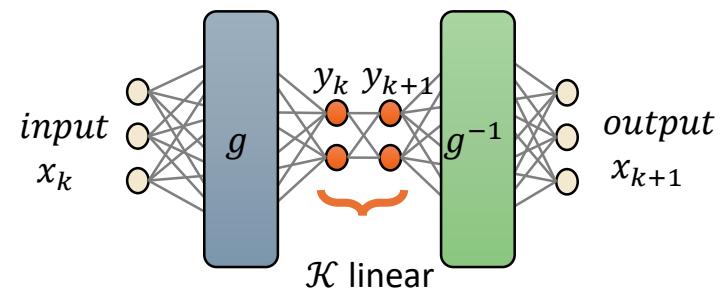
K-mixup: Data augmentation for offline reinforcement learning using mixup in a Koopman invariant subspace

Junwoo Jang ^{a,1}, Jungwoo Han ^{b,1}, Jinwhan Kim ^{c,*}

^aDepartment of Naval Architecture and Marine Engineering, University of Michigan, Ann Arbor, MI, 48109, USA

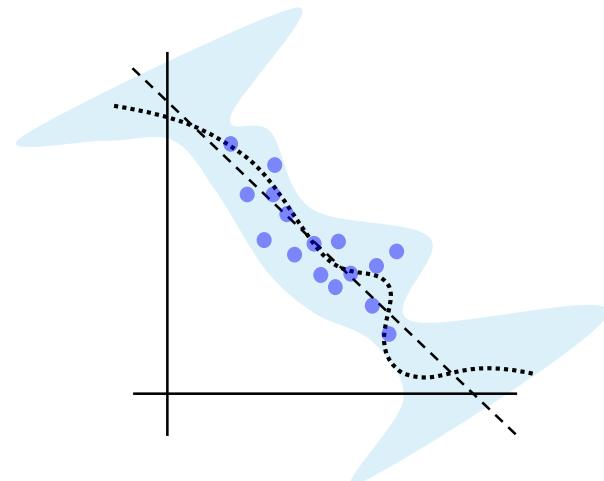
^bHyundai Motor Company, Seoul, 06182, South Korea

^cDepartment of Mechanical Engineering, KAIST, Daejeon, 34141, South Korea



Model learning with uncertainties (1), (2)

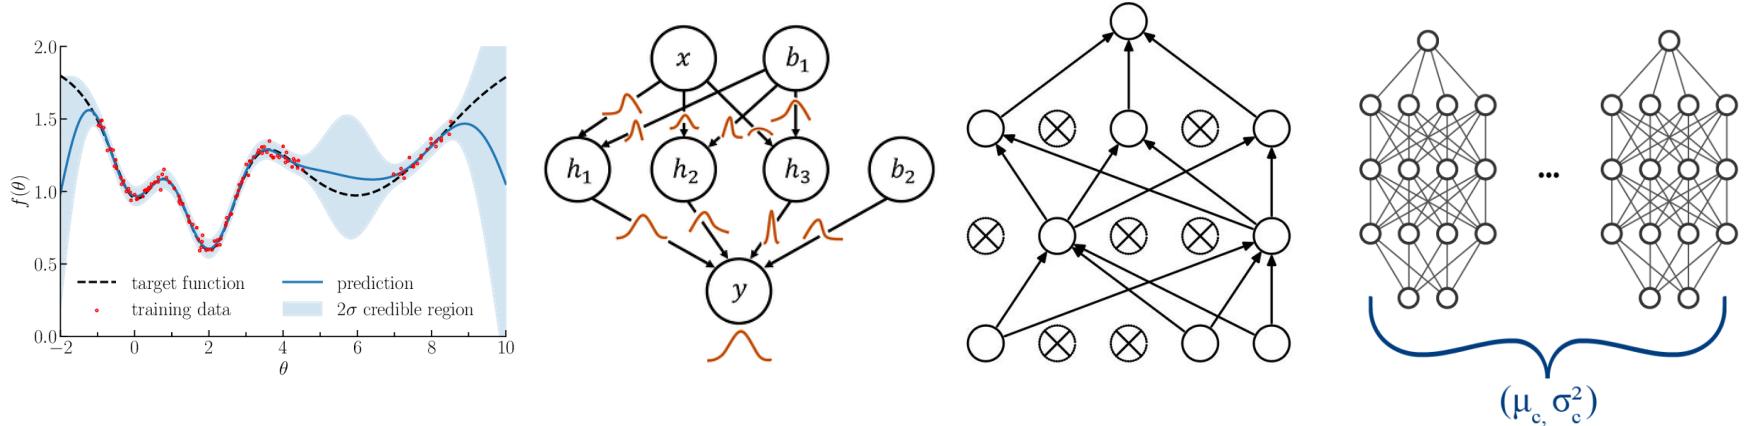
- > Make f_θ as a probabilistic model to predict uncertainties
 - discourage overconfidence by jointly learning the uncertainties
 - use the uncertainties for the control (and data collection)
 - distribution mismatch
 - training time: explore the uncertain space
 - test time: exploit the certain space



Model learning with uncertainties (1), (2)

> How to obtain the uncertainties?

- Gaussian process -> PILCO (2011)
- Bayesian neural networks
- Monte-Carlo dropout
- Model ensemble -> PETS (2018), MBPO (2019)
- Latent stochasticity -> Dreamer (2021)
- Evidential dynamics

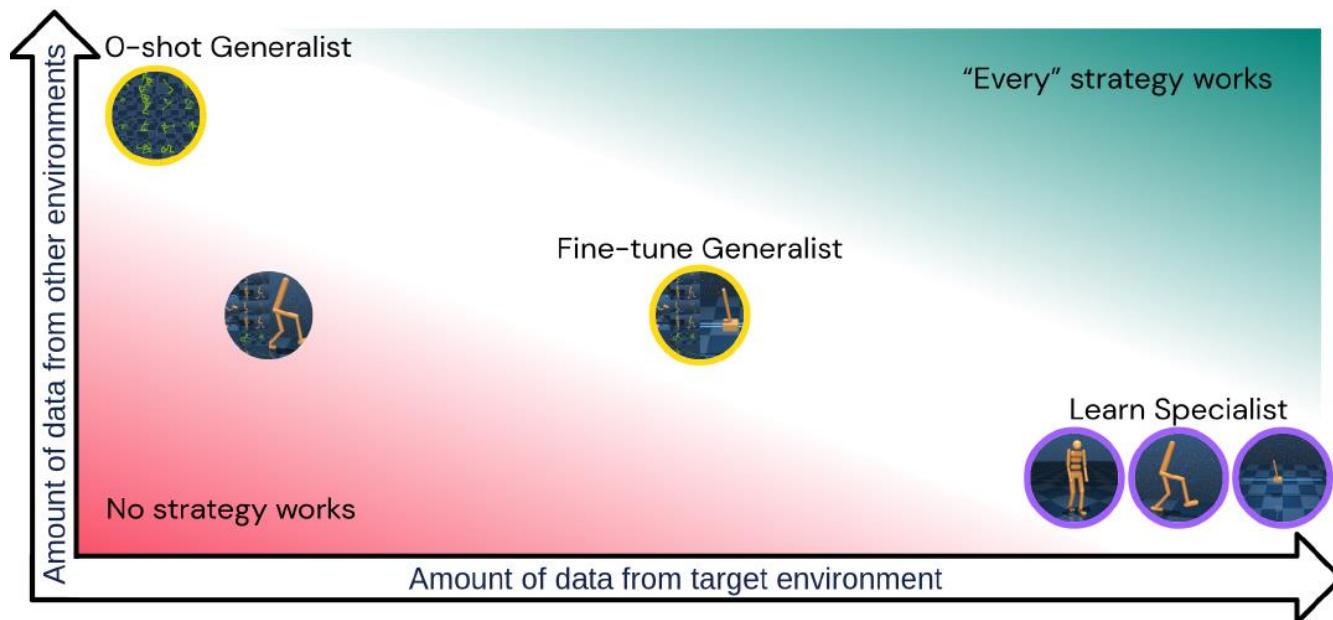
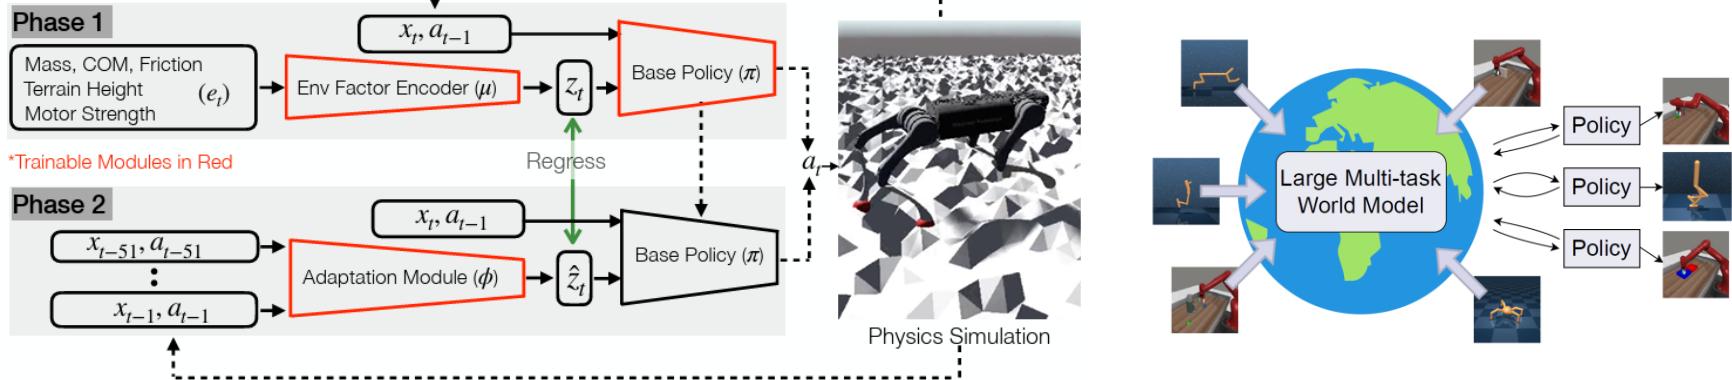


Meta-learning dynamics, foundation model (3)

- > Adapt model parameters in real time
 - meta-learning dynamics + real-time dynamics adaptation
 - meta-learning dynamics + adaptive control
 - policy conditioned on model parameters
- > Learn a generalized model (foundation model)
- > Examples
 - Learning to adapt in dynamic, real-world environments through meta-RL
 - Neural-fly enables rapid learning for agile flight in strong winds
 - RMA: rapid motor adaptation for legged robots
 - GNM: a general navigation model to drive any robot
 - Drive anywhere: Generalizable end-to-end autonomous driving with multi-modal foundation models
 - A generalist dynamics model for control

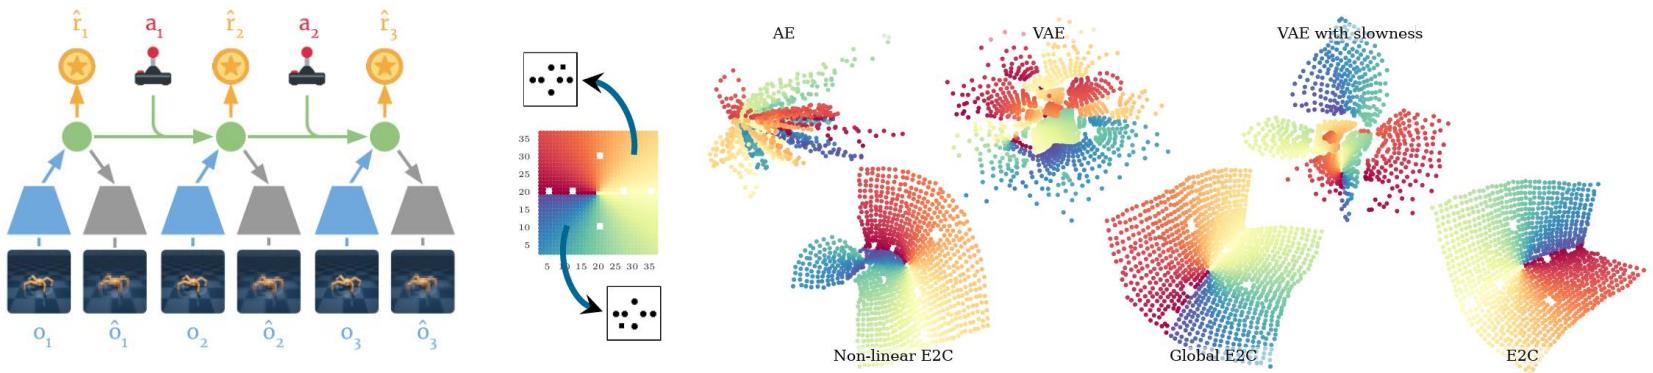
Meta-learning dynamics, foundation model (3)

A) Training in Simulation



Rich observations (4)

- > learning world model from high-dimensional sensory input (e.g., images)
 - Intersection between CV and MBRL
 - Representation/CV techniques (e.g., contrastive learning, diffusion)

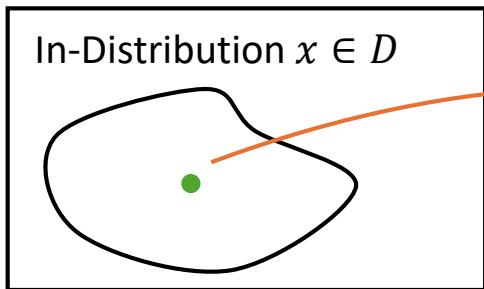


- > Examples
 - Embed to control: a locally linear latent dynamics model for control from raw images
 - Planet: learning latent dynamics for planning from pixels
 - dreamer v1, v2, v3
 - TD-MPC 1, 2

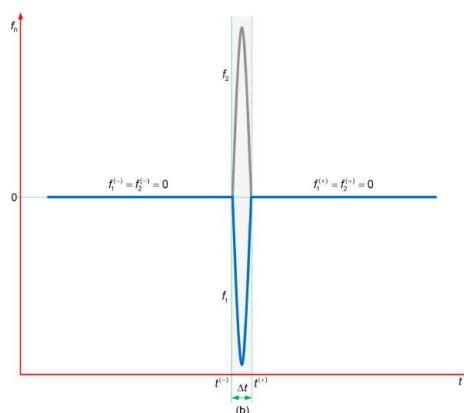
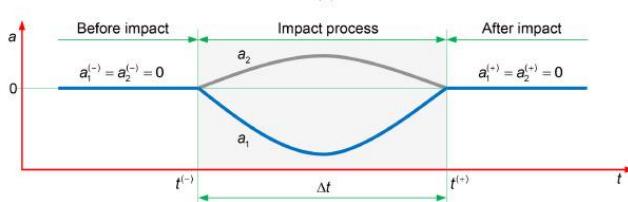
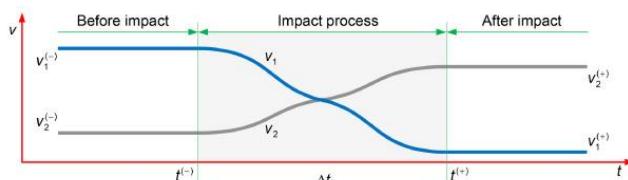
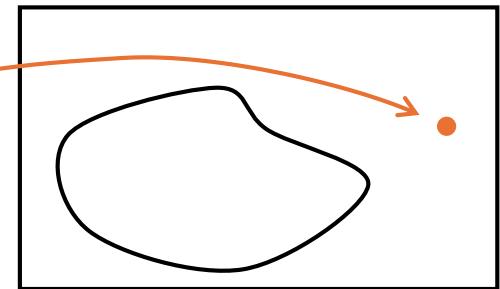
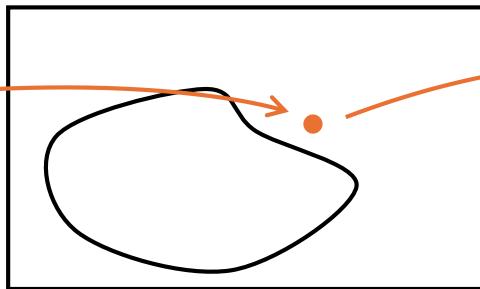
Underfitting and compounding errors (5)

- > Underfitting: hard to model complex/non-smooth dynamics (e.g., contacts)
- > Compounding error: in multi-step prediction (for planning), the iterative predictions produce large errors

Data space $x \in R^n$



Prediction must be always in distribution



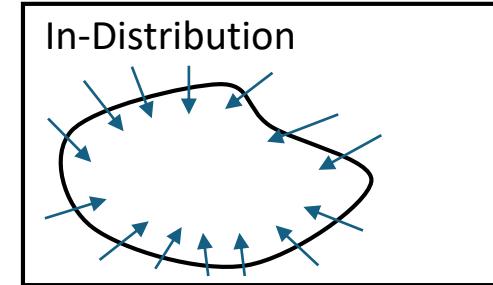
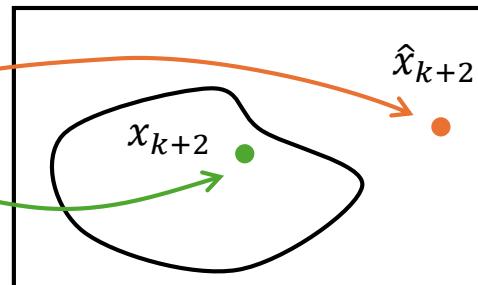
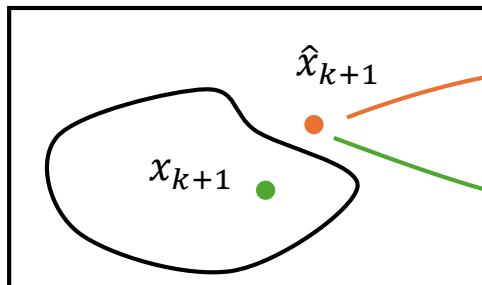
Underfitting and compounding errors (5)

> Solution

- Keep in-distribution (better representation)
- hierarchical or hybrid model
- curriculum learning
- fine-tuned by model-free RL

> Examples

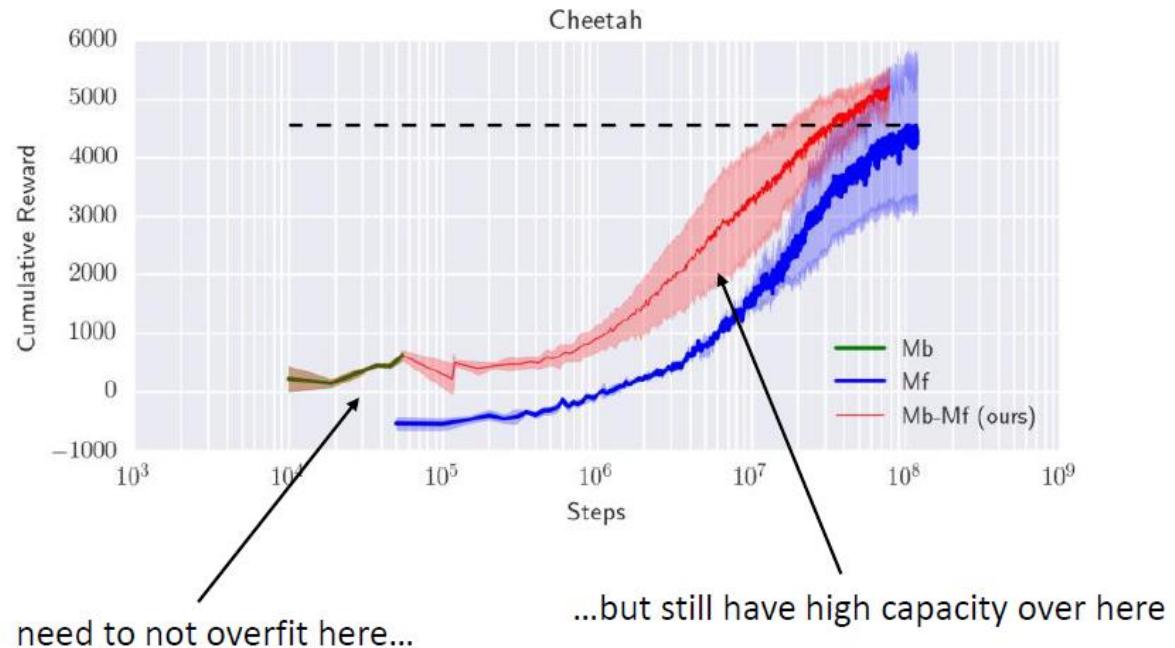
- latent dynamics for contact-rich control (hybrid latent model)
- hydrodynamic modeling of a robotic surface vehicle using representation learning for long-term prediction



Underfitting and compounding errors (5)

> Example

- Neural Network Dynamics for Model-Based Deep Reinforcement Learning with Model-Free Fine-Tuning



Model learning

- > Other kinds of models
 - so far: modeling $p(s_{t+1}|s_t, a_t)$ or $p(z_{t+1}|z_t, a_t)$ and $p(z_t|s_{t-n:t})$
- > Many alternatives
 - inverse model $p(a_t|s_t, s_{t+1})$
 - multi-step inverse model $p(a_{t:t+n}|s_t, s_{t+n})$
 - future prediction without actions $p(s_{t+1:t+n}|s_t)$
 - video interpolation $p(s_{t+1:t+n}|s_t, s_{t+n+1})$
 - transition distribution $p(s_t, a_t, s_{t+1})$

When to use model-based RL?

- > Big upsides and big downsides
 - immensely useful, far more data efficient if model is easy to learn
 - model can be trained on data without reward labels (fully self-supervised)
 - model is somewhat task-agnostic (can be transferred to other tasks)
 - model don't optimize for task performance
 - sometimes harder to learn than a policy
 - another thing to train, more hyperparameters, more compute intensive