

ECE7121 Learning-based control – 2025 Fall

# Inverse RL and Curriculum learning



INHA UNIVERSITY

## Inverse RL

- > Offline RL: learning from suboptimal or mixed demonstrations
  - reward labeling: at least, we know the relative goodness
- > Imitation learning: copy the good demonstrations
  - no reward labeling
- > Inverse RL: a type of learning from good demonstrations
  - similar to IL, no reward available
  - learn a reward function such that

$$\pi^* = \arg \max_{\pi \in \Pi} \mathbb{E}_{\pi, p}[r^*(s, a)]$$

- learn policy given the reward

# Inverse RL

> Why don't we directly learn a policy?


- assumes learning the reward function is statistically easier than directly learning the policy
- (perhaps) closer to how human learns



# Inverse RL

> Why don't we directly learn a policy?

- more interpretable (we can reason what an expert is trying to do)
- can generalize or adapt to new environments via RL
- manually design reward function can be difficult



**Table 1.** Reward terms for task, stability, and smoothness. We adapt the reward structure from Rudin et al. (2021) to our robot and task.

Term	Equation	Weight
$r_{v_x^{\text{cmd}}}$ : xy-vel	$\exp\{- \mathbf{v}_{xy} - \mathbf{v}_{xy}^{\text{cmd}} ^2 / \sigma_{vxy}\}$	0.02
$r_{\omega_z^{\text{cmd}}}$ : yaw-vel	$\exp\{-(\omega_z - \omega_z^{\text{cmd}})^2 / \sigma_{\omega z}\}$	0.01
z vel	$\mathbf{v}_z^2$	-0.04
roll-pitch vel	$ \boldsymbol{\omega}_{xy} ^2$	-0.001
height	$(h - h^0)^2$	-0.6
orientation	$ \mathbf{g}_{xy}^{\text{ori}} ^2$	-0.002
self-collision	$\mathbf{1}_{\text{selfcollision}}$	-0.02
joint limit	$\mathbf{1}_{g_i > g_{max}} \mathbf{1}_{g_i < g_{min}}$	-0.2
joint torques	$ \boldsymbol{\tau} ^2$	-2e-7
joint accel	$ \ddot{\mathbf{q}} ^2$	-5e-9
action rate	$ \mathbf{a}_{t-1} - \mathbf{a}_t ^2$	-2e-4
foot airtime	$\sum t_{air} * \mathbf{1}_{\text{new contact}}$	0.02

# Inverse RL

- > Why should we worry about learning rewards?

Computer Games  
reward



Mnih et al. '15

Real World Scenarios


robotics



dialog






autonomous driving



what is the reward?  
often use a proxy

- might take very different actions
- by itself, this is an underspecified problem  
(many reward function can explain the same behavior)



## Linear reward formulations

- > Commonly used in IRL algorithms
  - easier to optimize
  - $\theta$ : unknown weight,  $\phi$ : known features

$$R_\theta(s) = \theta_1\phi_1(s) + \dots + \theta_d\phi_d(s) = \boldsymbol{\theta}^\top \boldsymbol{\phi}(s)$$

- Why linear?
  - value function is linear in expected features
  - $V^\pi(s) = \mathbb{E}\left[\sum_{t=0}^{\infty} \gamma^t r(s^{(t)}) | s^{(0)} = s\right]$

$$= \boldsymbol{\theta}^\top \mathbb{E}\left[\sum_{t=0}^{\infty} \gamma^t \boldsymbol{\phi}(s^{(t)}) | s^{(0)} = s\right]$$

$$= \boldsymbol{\theta}^\top \mu^\pi(s)$$

# Linear reward formulations


- > Key idea: the expert is better than other policies

- other policies  $\pi_n$ : RL, manually designed, etc
- Linear program IRL (LP-IRL)

$$\hat{\theta} = \arg \max_{\theta} \{ \sum_i \theta^T (\mu^{\pi^*} - \mu^{\pi_i}) \} \text{ s.t. } |\theta_d| \leq 1 \forall d$$

- this process can be iterative: solve  $\theta$ , train new  $\pi_n$ , solve new  $\theta$ , ...
- Max margin IRL (apprenticeship learning)
  - adopt the max margin philosophy from SVM

$$\min_{\theta} \lambda \|\theta\|_2 + \sum_i \xi_i \text{ s.t. } \theta^T (\mu^{\pi^*} - \mu^{\pi_i}) \geq 1 - \xi_i$$



## Linear reward formulations

- > Max margin IRL (apprenticeship learning)
  - applications in autonomous helicopter aerobatics (2010, IJRR)
    - first, obtain reward function from expert demonstrations
    - obtain policy using Gauss-Newton LQR under real helicopter dynamics
    - shows in-place flips, in-place rolls, loops and hurricanes, auto-rotation landings, chaos and tic-tocs



# Linear reward formulations




## > Max entropy IRL (MaxEnt IRL)

- IRL is an ill-posed problem
  - many reward functions correspond to the same policy
  - many stochastic mixtures of policies correspond to the same feature expectation
- Maximum entropy principle
  - The probability distribution which best represents the current state of knowledge is the one with largest entropy
  - The distribution with largest policy (uncertainty) makes the least assumptions about the true distribution
  - $\max_p \{-\sum p(\tau) \log p(\tau)\}$

# Linear reward formulations

## > Max entropy IRL (MaxEnt IRL)

- Given known constraints (prior data), choose the distribution that satisfies them while maximizing entropy
  - unknown parts are random → most random is the largest entropy
  - as random as possible while matching features
  - ex) suppose the probability of rain tomorrow is 30%, and we know nothing else
  - ex) IRL example: maximize the entropy of trajectory distribution  $p(\tau)$



- Given the reward function, obtain the stochastic (soft-optimal) policy via the soft Bellman updates, then match feature expectations

## Linear reward formulations

### > Max entropy IRL (MaxEnt IRL)

- Maximizing the entropy → maximize the likelihood of the observed data under the maximum entropy distribution
- from trajectories to states: successful imitation boils down to learning a policy that matches the state (or state/action) visitation distribution
- Algorithm (with known dynamics, linear costs)

0. Initialize  $\theta$ , gather demonstrations  $\mathcal{D}$
1. Solve for optimal policy  $\pi(a|s)$  w.r.t.  $c_\theta$  with value iteration
2. Solve for state visitation frequencies  $p(s | \theta, T)$
3. Compute gradient  $\nabla_\theta \mathcal{L} = \frac{1}{M} \sum_{\tau_d \in \mathcal{D}} \mathbf{f}_{\tau_d} - \sum_s p(s | \theta, T) \mathbf{f}_s$
4. Update  $\theta$  with one gradient step using  $\nabla_\theta \mathcal{L}$




## Linear reward formulations


- > MaxEnt IRL requires
  - solving for (soft) optimal policy in the inner loop
  - enumerating all state-action tuples for visitation frequency and gradient
  - reward was assumed linear over features  $\phi$ 
    - general function approximations for the reward
- > To apply this in practical problem settings, we need to handle
  - large and continuous state and action spaces
  - states obtained via sampling only
  - unknown dynamics
    - sampled based approximations for the partition function

# GAN framework

- > IRL framework




- > Guided cost learning (2016 ICML)



# GAN framework

- > Generative adversarial network (GAN)
  - train generator and discriminator iteratively



- > Inverse RL as a GAN
  - policy changed to make it harder to distinguish from demos




# GAN framework

- > Generative adversarial imitation learning (GAIL, 2016)
  - discriminator learns to tell expert state action pairs from the agent's
  - end-to-end imitation without hand-crafted rewards
  - often simpler to set up optimization


Guided Cost Learning

Finn et al., ICML 2016




Generative Adversarial Imitation Learning

Ho & Ermon, NIPS 2016




## Generalization via inverse RL

- > What can we learn from the demo to enable better transfer?
  - need to decouple the goal from the dynamics
  - policy = reward + dynamics



demos



reproduce behavior  
under different conditions

# Goal classifier

- > Reward can be explicitly trained if we know the objective
  - learn to discern goal states from other states

Example task: put pencil case behind notebook


Positive examples



Negative examples



Trained binary classifier




Use output as reward signal.

- collect examples of successful / unsuccessful states
- train binary classifier
- run RL with classifier as reward

# Goal classifier

- > Reward can be explicitly trained if we know the objective
  - RL algorithm will seek out states that the classifier thinks are good
  - It may simply find states that the classifier wasn't trained on
    - exploiting the classifiers weakness
  - Add states that RL visits as negative examples for the classifier
    - classifier can't be exploited (update the classifier during RL)
    - what if some of the visited states are successful?
    - important to regularize the classifier (or balanced training datasets)



# Goal classifier

- > Goal classifiers for robotic RL example (2023 CoRL)
  - collect 50 demos, use final states as success state examples
  - initialize RL replay buffer with demos

BC (26/12% success rate)



Cloth hanging

MEDAL++ (62/46% success rate)



Bowl cloth cover



## Goal classifier

- > Pros and cons
  - practical framework for task specification
  - adversarial training can be unstable  
(though variety of regularization tricks from GAN literature)
  - requires examples of desired behavior or outcomes


- > Can humans provide feedback on policy roll-outs?
  - instead of requiring demos or example goals

A couple options



“How good is this trajectory?”



“Which trajectory is better?”




Relative preferences are easier to provide!



# RLHF

## > RL from human feedback (RLHF)


- often easier for people to make than handwriting a reward function
- often easier than providing scalar reward  
(how much do you like this ad?)
- needed 900 bits of feedback from a human evaluator to learn to backflip



# RLHF

## > RL with AI feedback (RLAIF)

- ask another language model ‘which of these responses is less harmful?’




## > Why human feedback matters in control

- Unmodeled objectives: comfort (low jerk), noise, safety margins, aesthetics—hard to hand-code in reward functions
  - social navigation
- Sim-to-real gaps: policies that track in sim can feel unsafe/annoying in reality
- Contextual trade-offs: users/environments change the weighting of goals; human preferences adapt the objective online.



# Curriculum learning


- > Training an agent with tasks of gradually increasing difficulty
  - simpler skills are mastered first and then built upon to solve complex tasks



- ex) learn how to walk (safe and easy env. → dangerous env.)





- ex) evolution of life on earth




# Curriculum learning

- > The importance of starting small (1993)
  - Humans display an exceptional capacity to learn
  - Humans are remarkable for the unusually long time it takes to reach maturity
  - Through culture, learning has created the basis for a non-genetically based transmission of behaviors which may accelerate the evolution of our species
- > Curriculum in supervised learning (2009)
  - clarify when and why a curriculum strategy can benefit



# Curriculum learning


- > Several types of curriculum learning
  - a special form of transfer learning where the initial tasks are used to guide the learner so that it will perform better on the final task




# Curriculum learning

## > Curricula for robot tasks


- exposing the learning agent to a sequence of tasks of increasing difficulty
  - increasing the number of moving joints
  - increasing speeds
  - using easier initial robot configurations
  - change the cost function (increase the coefficient on the torque cost)



(a) Sub-task 1



(b) Sub-task 2






(c) Sub-task 3

# Curriculum learning


## > Curricula for robot tasks

- change the physics (meta-learning)
- change the environment




# Curriculum learning

- > How to automatically generate a curriculum for a given RL agent
  - given metrics of the agent's behavior like performance or visited states, automatic curriculum learning methods generate new tasks
    - initial/target state: start learning from states close to a given target (sparse reward)
    - reward function: curiosity-based approaches, skill discovery
    - environments / opponent (for self-play algorithms)



# Curriculum learning

- > Ex) Automatic domain randomization (2019)
  - Environment parameters are sampled from independent uniform dist.
  - If agent performance > upper threshold → expand parameter ranges
  - If performance < lower threshold → contract ranges



- still, there are tuning parameters (update step size)
  - how much to expand/contract the sampling distribution?

## Curriculum learning

- > Ex) self-paced deep RL (2020)
  - ADR relies only on success thresholds
    - need a framework that adapts smoothly and theoretically grounded
  - Key idea: learn a task/context distribution

$$\max_{v,w} J(v, w) - \alpha D_{KL}[p_v(c) \parallel \mu(c)]$$

$p_v(c)$  is a current task distribution,  $\mu(c)$  is a target distribution

- Start with a simple task distribution (easy goals, stable conditions)
- Evaluate the value function → shift distribution toward tasks that improve learning
- Use KL divergence to make the current distribution close to the target distribution gradually

## Curriculum learning

- > Legged robot walking policy



## Reference

- > Curriculum learning
  - <https://lilianweng.github.io/posts/2020-01-29-curriculum-rl/>